
FORGE: REACTION-GROUNDED GENERATIVE DESIGN OF IONIZABLE LIPIDS

Anonymous authors

Paper under double-blind review

ABSTRACT

Ionizable lipids are central to RNA delivery, yet their discovery still largely relies on modular synthesis and screening of molecular libraries whose constituent chemistry is chosen in advance. De novo generation could expand this design space, but opening molecular identity risks severing the reaction semantics that make a design experimentally interpretable and assemblable. We introduce FORGE (**F**low-matched, **O**pen-ended, **R**oute-resolved **G**eneration and **E**xploration), a reaction-grounded generative framework that uses known reaction chemistry as the coordinate system for molecular generation. Reaction programs specify precursor roles, reaction-core positions, transformation order, and coarse molecular architecture, defining how regions of the final molecule participate in synthesis while leaving their molecular identities open. A conditional discrete flow then generates the complete molecular realization at atom-and-bond resolution without component identifiers, stored precursor graphs, or fragment tokens. Across Ugi, repeated aza-Michael addition, and repeated reductive amination, FORGE achieves exact level-one transform-consistent yields of 78.7%, 77.5%, and 31.3%, substantially outperforming matched whole-graph generation followed by post-hoc reaction filtering. Across all three programs, FORGE also generated exact-L1 products with recovered precursor identities beyond the corresponding training component catalogues, a capability unavailable to finite catalogue assemblers by construction. FORGE therefore uses reaction knowledge to preserve assembly semantics while opening molecular identity to de novo generation.

1 INTRODUCTION

Lipid nanoparticles (LNPs) deliver nucleic-acid therapeutics including small interfering RNA, mRNA vaccines and *in vivo* genome-editing reagents (Adams et al., 2018; Polack et al., 2020; Baden et al., 2021; Gillmore et al., 2021). Within each LNP, the ionizable lipid supports nucleic-acid complexation and influences particle assembly and endosomal escape, making its structure a central determinant of function (Cullis & Hope, 2017; Hou et al., 2021). Small changes to headgroup and hydrophobic structure can substantially alter delivery, yet their effects depend on the complete molecular context and do not consistently follow simple design rules (Cheng et al., 2020; Han et al., 2021; Hatit et al., 2023). Discovery therefore commonly proceeds by defining modular precursor libraries, synthesizing combinations and screening the resulting LNPs experimentally (Akinc et al., 2008; Love et al., 2010; Li et al., 2024; Xu et al., 2024). This workflow preserves experimental tractability, but largely fixes the molecular vocabulary of the search before design begins.

De novo generative modeling offers a way beyond this ceiling, but reaction grounding creates a representational problem. Synthesis-aware generators commonly encode chemical feasibility by constructing molecules through predefined reactions and reactant inventories (Gao et al., 2022; Koziarski et al., 2024; Cretu et al., 2025; Rekesh et al., 2026). This makes chemistry part of the generative action space and preserves an explicit assembly path, but also ties reachable molecular identity to the supplied vocabulary. Generic whole-molecule generators have the complementary strength: they can generate new molecular identities, but do not by themselves preserve which regions must function as particular precursors or which

transformations should assemble them (Qin et al., 2025; Lee et al., 2025). A complementary line of work evaluates synthesizability after generation rather than encoding construction directly into the generative state (Gao & Coley, 2020). Recent work has also begun to explore de novo generative modeling specifically for ionizable lipids (Ou et al., 2024; Maganti et al., 2026). We ask whether known reaction chemistry can instead organize an open generative state without specifying the molecular structures that occupy it. Our premise is that reaction grounding need not require component enumeration: reaction knowledge can provide semantic coordinates that determine how molecular regions participate in construction while leaving their identities generative.

We introduce **FORGE** (**F**low-matched, **O**pen-ended, **R**oute-resolved **G**eneration and **E**xploration), a reaction-grounded generative framework that uses known reaction chemistry as the coordinate system for molecular generation. A reaction program specifies the assembly chemistry, precursor roles, reaction-core positions, transformation order and coarse molecular architecture, but not precursor identities. Conditioned on this program, a discrete flow generates the complete molecular realization at atom-and-bond resolution without component identifiers, stored precursor graphs or fragment tokens. Because reaction semantics are preserved in the generative state, each completed molecule can be decomposed into the precursors implied by its program and the specified assembly reaction replayed exactly. Precursors that are not directly available are then assessed recursively for upstream preparation, yielding a complete evidence-qualified computational synthesis dossier or an explicit unresolved or search-censored outcome. We evaluate this formulation across a Ugi three-component reaction, repeated aza-Michael addition and repeated reductive amination. These experiments test generalization across specified reaction programs rather than zero-shot support for arbitrary chemistry. Our contributions are threefold. First, we formulate reaction-grounded whole-graph generation in which reaction programs provide semantic coordinates while molecular identity remains generative at atom-and-bond resolution. Second, matched controls across three assembly programs show that reaction semantics substantially increase exact transform-consistent generation relative to whole-graph generation followed by post-hoc filtering. Third, FORGE leaves precursor identity open rather than restricting generation to a finite component vocabulary; across all three programs, this yields exact-assembly products with recovered precursor identities beyond the corresponding training catalogues.

2 FORGE

We formulate ionizable-lipid design as conditional generation under a specified reaction program, while leaving precursor identities unspecified. Let $G \in \mathcal{G}$ denote a complete ionizable-lipid graph and $P \in \mathcal{P}$ a reaction program specifying the assembly chemistry, precursor roles, reaction-core positions, and transformation order. FORGE models the conditional distribution $p_\theta(G | P)$ without access to precursor identifiers, stored component graphs, or fragment tokens. Given P , FORGE generates a complete molecular graph G . A downstream bounded program constructor then returns a recursive synthesis program T or an explicit abstention outcome. Writing $\mathfrak{T}$ for the space of recursive synthesis programs and $\perp$ for that abstention,

$$G \sim p_\theta(\cdot | P), \quad T = \mathcal{A}(G, P) \in \mathfrak{T} \cup \{\perp\}, \quad (1)$$

where $\mathcal{A}$ is the bounded program constructor that verifies the final assembly, searches for preparation routes to unavailable precursors, and checks terminal supplier evidence. P and $\mathcal{A}$ act at different stages: P guides the generation of G , whereas $\mathcal{A}$ assesses the completed graph without modifying it. A pair (G, T) is designated route-complete only when G passes molecular and exact-assembly verification and every required precursor traces to an available starting material; otherwise the route assessor abstains.

A reaction program is a tuple $P = (c, d, \mathbf{r}, \mathbf{k}, M)$, where c identifies the ordered assembly chemistry, d is the program depth, $\mathbf{r}$ assigns precursor roles, $\mathbf{k}$ marks reaction-core positions, and M specifies coarse per-role morphology as four integers per role: exterior atom count, junction budget, cycle rank and attachment count. M sets only coarse structural constraints; the model generates the exact topology, atoms and bonds of each region. FORGE samples




Figure 1: **FORGE uses reaction chemistry as the coordinate system for molecular generation.** A reaction program specifies precursor roles and reaction-core structure without fixing molecular identities; a conditional discrete flow generates the complete lipid, and exact assembly verification requires forward reconstruction from the implied precursors. Optional downstream routing assesses unavailable precursors or abstains. Bottom, FORGE combines explicit reaction semantics with generated molecular identity. Colors denote semantic roles, not component identities.

M from a training-fold morphology prior and then generates G conditioned on the resulting program, so that $p_\theta(G, M | c) = p_{\text{prog}}(M | c) p_\theta(G | M, c)$. For each graph position i , FORGE combines learned embeddings of program identity, program depth, precursor role and reaction-core position, $h_i^P = \text{LN}(e_c(c) + e_d(d) + e_r(r_i) + e_k(k_i))$. These coordinates tell the generator which precursor role each position plays and whether it lies in the reaction core, but not which molecular structure should occupy that role.

The framework comprises three core stages followed by optional downstream route assessment: the program specifies roles, core positions, order and coarse morphology; a conditional discrete flow generates every atom, bond and topology state under it; a reverse transform decomposes the product and the corresponding forward transform must reconstruct it exactly; and precursors that are not already available may then be assessed recursively, with abstention when the route evidence is incomplete. Generation operates on a sparse tree-plus-closure encoding. Let $X_G = (\mathbf{a}, \boldsymbol{\pi}, \mathbf{b}, \mathbf{e}^L, \mathbf{e}^R, \mathbf{z})$ be the categorical graph state, with $\mathbf{a}$ atom states, $\boldsymbol{\pi}$ parent pointers, $\mathbf{b}$ spanning-tree bond states, $\mathbf{e}^L$ and $\mathbf{e}^R$ closure endpoints and $\mathbf{z}$ closure-bond states. The decoder is surjective onto the declared graph space $\mathfrak{G}_{N_{\text{max}}, K}$, here $N_{\text{max}} = 194$ and $K = 3$, so every graph within the declared atom, bond and size support is representable. This is a statement about architectural support, not about reaction consistency or probability under a trained model.

The chemistry adapter marks the atom, pointer and bond states that P fixes. FORGE holds those constant while corrupting the remaining states along the linear probability path and predicting their clean values from (X_t, t, P) . Writing $\mathcal{Q}$ for the six state families, $m_j^{(q)}(P) \in \{0, 1\}$ for whether coordinate j remains variable under P , and μ_{bal} for the frozen program-balanced training distribution, and $\ell_j^{(q)}(\theta) = -\log p_\theta(X_{G,j}^{(q)} | X_t, t, P)$ for the per-

Table 1: **One shared conditional flow learns three ionizable-lipid assembly programs.** Exact-L1 yield and decomposition/replay statistics are reported as mean $\pm$ sample standard deviation across three independent training seeds. Results are never pooled across programs. Per-seed results and full matched controls are reported in Appendix C.

Program	FORGE conditioned	Shared null	Cyclic program	Decomposition / replay (%)
Ugi	78.7 $\pm$ 10.1	11.4 $\pm$ 1.1	40.2 $\pm$ 0.7	90.7/100.0
Aza-Michael	77.5 $\pm$ 3.7	4.0 $\pm$ 0.8	3.8 $\pm$ 0.6	100.0/100.0
Reductive amination	31.3 $\pm$ 14.3	13.2 $\pm$ 1.3	18.1 $\pm$ 0.8	84.3/100.0

coordinate denoising term, FORGE minimizes

$$\mathcal{L}_{\text{FORGE}}(\theta) = \mathbb{E}_{(G,P) \sim \mu_{\text{bal}}, t \sim \mathcal{U}(0,1), X_t \sim p_t(\cdot | X_G, P)} \left[\sum_{q \in \mathcal{Q}} \lambda_q \frac{\sum_j m_j^{(q)}(P) \ell_j^{(q)}(\theta)}{\max\{1, \sum_j m_j^{(q)}(P)\}} \right], \quad (2)$$

where λ_q weights family q and a family with no variable coordinates contributes zero loss. The corruption and sampling operators project every intermediate state onto the program-fixed coordinates, so those coordinates are invariant by construction while the masked objective learns the conditional clean-state distributions of the remainder. This is the precise sense in which reaction semantics enter generation rather than filtering: the program fixes which coordinates the model is never permitted to resample. Appendix B gives the sampler, verifier and bounded route assessor in full, and Appendix E the representation, invariance, sampling-efficiency and termination results.

3 RESULTS

We evaluate three assembly programs, with Ugi as the deeply supervised case, repeated aza-Michael addition as a second family, and repeated reductive amination as a lighter stress test, across four matched training arms: Ugi-only conditioning, shared conditioning, shared training without program information, and shared training with cyclically incorrect program labels. All arms share architecture, optimization budget and attempts per program, and we evaluate the prespecified final checkpoint from each of three independent seeds. The primary metric is exact-L1 yield per generation attempt: generated graphs passing molecular validity, program-aware decomposition and exact forward replay, divided by total attempts. Training uses only the component-disjoint training folds, and held-out products neither enter training nor select checkpoints. Generated molecules are not treated as independent replicates. Appendix C states the full protocol, including the external-baseline admission rules.

3.1 ONE SHARED MODEL LEARNS MULTIPLE ASSEMBLY PROGRAMS

Across 3,072 held-out generation attempts per program in each of three independent training seeds, the final conditioned model achieved exact-L1 yields of 78.7 $\pm$ 10.1% for Ugi (72.3%, 90.3% and 73.5% by seed), 77.5 $\pm$ 3.7% for repeated aza-Michael addition (76.5%, 81.5% and 74.3%), and 31.3 $\pm$ 14.3% for repeated reductive amination (38.6%, 14.8% and 40.4%) (mean $\pm$ sample standard deviation; Table 1). Exact forward-replay precision was 100% in every program and seed, but mean decomposition ambiguity was 54.9% for aza-Michael addition and 74.1% for reductive amination, and the latter varied substantially between seeds. Relative to Ugi-only training (25.4 $\pm$ 12.0%), shared conditioning increased Ugi exact-L1 yield to 78.7 $\pm$ 10.1%, a 53.26-point improvement (95% paired-seed interval 41.24 to 61.82), met every prespecified Ugi-retention gate, and produced zero fixed-state violations and zero support overflows across 27,648 held-out attempts. One conditional flow can therefore learn the three specified reaction-program coordinates without erasing the deeply supported Ugi behavior.

Table 2: **FORGE combines exact Ugi assembly with precursor identities beyond a finite training vocabulary.** All methods are evaluated under the common component-disjoint split, three-seed attempt budget and exact-L1 verifier, as mean $\pm$ sample standard deviation per 1,000 generation attempts. Open-ended yield denotes exact-L1 products with a verifier-returned precursor constitution absent from the method-visible training catalogue; the stricter yield against the full training catalogue is reported in Table 6. $\dagger$ denotes a structural zero for methods restricted to finite component inventories. N/E, not estimable. No published values from the original baseline papers are inserted. Full admission rules and seed-level results are reported in Appendix C.

Method	Valid	Exact L1	Unique L1	Open-ended	Held-comp.	Diversity
Reaction-space external baselines						
RGFN	0.0 $\pm$ 0.0	0.0 $\pm$ 0.0	0.0 $\pm$ 0.0	0.0 $\pm$ 0.0 †	0.0 $\pm$ 0.0	N/E
Generic whole-molecule external baselines with common post-hoc assessment						
DeFoG unconditional	346.5 $\pm$ 42.3	280.6 $\pm$ 26.6	280.6 $\pm$ 26.6	280.6 $\pm$ 26.6	15.5 $\pm$ 2.4	0.727 $\pm$ 0.014
GenMol/SAFE	627.5 $\pm$ 6.2	0.0 $\pm$ 0.0	0.0 $\pm$ 0.0	0.0 $\pm$ 0.0	0.0 $\pm$ 0.0	0.873 $\pm$ 0.009
Matched controls and formulation baselines						
Learned inventory selector	1000.0 $\pm$ 0.0	1000.0 $\pm$ 0.0	849.5 $\pm$ 13.7	0.0 $\pm$ 0.0 †	0.0 $\pm$ 0.0	0.717 $\pm$ 0.001
Finite catalogue oracle	1000.0 $\pm$ 0.0	1000.0 $\pm$ 0.0	837.6 $\pm$ 8.2	0.0 $\pm$ 0.0 †	0.0 $\pm$ 0.0	0.719 $\pm$ 0.000
Shared-null + post-hoc	194.7 $\pm$ 15.6	114.4 $\pm$ 10.9	108.5 $\pm$ 9.7	108.5 $\pm$ 9.7	0.4 $\pm$ 0.2	0.754 $\pm$ 0.014
FACT-matched	60.8 $\pm$ 20.2	52.0 $\pm$ 15.6	50.8 $\pm$ 14.4	50.8 $\pm$ 14.4	0.0 $\pm$ 0.0	0.667 $\pm$ 0.009
FACT-generous	500.7 $\pm$ 213.8	296.7 $\pm$ 186.0	288.3 $\pm$ 175.3	288.3 $\pm$ 175.3	0.1 $\pm$ 0.2	0.774 $\pm$ 0.007
FORGE Transformer	866.0 $\pm$ 90.2	786.8 $\pm$ 100.8	751.2 $\pm$ 80.2	751.2 $\pm$ 80.2	0.1 $\pm$ 0.2	0.621 $\pm$ 0.011

3.2 REACTION SEMANTICS ALTER GENERATION RATHER THAN MERELY FILTERING ITS OUTPUT

Under matched attempt budgets, conditioned generation produced 786.8, 774.6 and 312.7 exact-L1 products per 1,000 attempts for Ugi, repeated aza-Michael and repeated reductive amination, compared with 114.4, 39.7 and 131.7 for matched null generation followed by post-hoc filtering. The paired conditioned-minus-post-hoc differences were 67.24 percentage points for Ugi (95% interval 60.84 to 79.95), 73.49 for aza-Michael (69.69 to 78.45), and 18.10 for reductive amination (2.51 to 26.04). A cyclic control that retained precursor roles, reaction-core coordinates and program depth but assigned the wrong reaction identity changed exact-L1 yield by 38.43 percentage points for Ugi (32.00 to 49.41), 73.67 for aza-Michael and 13.15 for reductive amination. A separate fixed-network inference-time intervention similarly showed that nulling the program coordinates sharply reduced Ugi exact-L1 yield, confirming that the semantic coordinates affect generation rather than only downstream assessment (Appendix C.3).

3.3 WHOLE-GRAPH GENERATION PRODUCES PRECURSOR IDENTITIES BEYOND A FINITE VOCABULARY

Table 2 compares FORGE with external synthesis-aware generators, matched controls and a finite-catalogue oracle under one common Ugi benchmark, normalizing every yield by generation attempts. The learned finite-inventory selector reached 1000.0 $\pm$ 0.0 exact-L1 products per 1,000 attempts with zero open-ended component generation, imposed by its finite action space, whereas FORGE produced 786.8 $\pm$ 100.8 exact-L1 products and 751.2 $\pm$ 80.2 unique exact-L1 products per 1,000 attempts. All of the latter were open-ended against the method-visible training catalogue; against the full training catalogue the open-ended yield was 729.3 per 1,000 (Table 6). Among open-ended external generators, unconditional DeFoG produced 280.6 $\pm$ 26.6 exact-L1 products and recovered held components at 15.5 $\pm$ 2.4 per 1,000 attempts against 0.1 $\pm$ 0.2 for FORGE; GenMol/SAFE produced 627.5 $\pm$ 6.2 valid connected outputs but no exact Ugi-L1 product, and RGFN emitted no admissible output under the strict native contract. Against the train-only assembler, FORGE effective component counts were 320.8, 156.9 and 102.9 versus 86.4, 23.4 and 16.9. Reaction consistency, open-ended identity generation and recovery of reserved components therefore trade off separately.




Figure 2: **Deterministic examples of open-ended FORGE ionizable-lipid graphs.** Each row is one seed-0 Ugi product with a unique exact-L1 decomposition and a recovered precursor constitution outside the method-visible training inventory. Left, the graph sampled jointly by FORGE; center, precursor origins recovered by exact atom-mapped forward replay, halos marking the Ugi core; right, the exact precursor constitutions returned by the verifier. Colors are semantic coordinates, not a vocabulary the model selects from. These are display examples, not experimental candidates, and exact L1 replay is a structural diagnostic, not evidence of activity or synthesis success. Selection rule and twelve further examples: Appendix F.

Accepted upstream transformations must reconstruct their targets under the corresponding forward reaction, and terminal branches must close under the frozen material-evidence policy. Because route coverage depends on the available reaction and procurement evidence, unresolved outputs are treated as abstentions rather than evidence of unsynthesizability; full route-assessment details are given in Appendix C.

4 DISCUSSION

FORGE introduces a reaction-grounded formulation of molecular generation in which known assembly chemistry provides semantic structure without fixing the molecular identities that occupy it. Across Ugi, repeated aza-Michael addition and repeated reductive amination, reaction-program conditioning substantially increased exact transform-consistent generation relative to matched whole-graph generation followed by post-hoc filtering. Whole-graph generation further produced exact assemblies with precursor identities beyond finite training catalogues, showing that reaction grounding need not reduce molecular design to component enumeration.

These results suggest a broader role for synthesis knowledge in generative modeling. Rather than entering only as a finite reaction action space or a terminal feasibility filter, known chemistry can define coordinates over an otherwise open molecular generative state. The evidence is strongest for Ugi, while repeated reductive amination remains substantially more variable and decomposition-ambiguous. The present study is computational, covers three specified reaction programs within bounded graph support, and treats recursive route assessment as a downstream evidence layer rather than experimental synthesis validation. Several prespecified outcomes did not favor FORGE and are reported in full in Appendix C. Within these limits, FORGE provides a practical route from fixed reaction semantics to open molecular identity.

REFERENCES

- David Adams, Alejandra Gonzalez-Duarte, William D. O’Riordan, Chih-Chao Yang, Mitsuharu Ueda, Arnt V. Kristen, Ivailo Tournev, Hartmut H. Schmidt, Teresa Coelho, John L. Berk, Kon-Ping Lin, Giuseppe Vita, Shahram Attarian, Violaine Planté-Bordeneuve, Michelle M. Mezei, Josep M. Campistol, Juan Buades, Thomas H. Brannagan, Byoung J. Kim, Jeeyoung Oh, Yesim Parman, Yoshiki Sekijima, Philip N. Hawkins, Scott D. Solomon, Michael Polydefkis, Peter J. Dyck, Pritesh J. Gandhi, Sunita Goyal, Jihong Chen, Andrew L. Strahs, Saraswathy V. Nochur, Marianne T. Sweetser, Pushkal P. Garg, Akshay K. Vaishnaw, Jared A. Gollob, and Ole B. Suhr. Patisiran, an RNAi therapeutic, for hereditary transthyretin amyloidosis. *New England Journal of Medicine*, 379(1):11–21, 2018. doi: 10.1056/NEJMoa1716153.
- Akin Akinc, Andreas Zumbuehl, Michael Goldberg, Elizaveta S. Leshchiner, Valentina Busini, Naushad Hossain, Sergio A. Bacallado, David N. Nguyen, Jason Fuller, Rene Alvarez, Anna Borodovsky, Todd Borland, Rainer Constien, Antonin de Fougères, J. Robert Dorkin, K. Narayanannair Jayaprakash, Muthusamy Jayaraman, Matthias John, Victor Kotliansky, Muthiah Manoharan, Lubomir Nechev, June Qin, Timothy Racie, Denitza Raitcheva, Kallanthottathil G. Rajeev, Dinah W. Y. Sah, Jürgen Soutschek, Ivanka Toudjarska, Hans-Peter Vornlocher, Tracy S. Zimmermann, Robert Langer, and Daniel G. Anderson. A combinatorial library of lipid-like materials for delivery of RNAi therapeutics. *Nature Biotechnology*, 26(5):561–569, 2008. doi: 10.1038/nbt1402.
- Lindsey R. Baden, Hana M. El Sahly, Brandon Essink, Karen Kotloff, Sharon Frey, Rick Novak, David Diemert, Stephen A. Spector, Nadine Rouphael, C. Buddy Creech, John McGettigan, Shishir Khetan, Nathan Segall, Joel Solis, Adam Brosz, Carlos Fierro, Howard Schwartz, Kathleen Neuzil, Lawrence Corey, Peter Gilbert, Holly Janes, Dean Follmann, Mary Marovich, John Mascola, Laura Polakowski, Julie Ledgerwood, Barney S. Graham, Hamilton Bennett, Rolando Pajon, Conor Knightly, Brett Leav, Weiping Deng, Honghong Zhou, Shu Han, Melanie Ivarsson, Jacqueline Miller, and Tal Zaks. Efficacy and safety of the mRNA-1273 SARS-CoV-2 vaccine. *New England Journal of Medicine*, 384(5):403–416, 2021. doi: 10.1056/NEJMoa2035389.
- Qiang Cheng, Tuo Wei, Lukas Farbiak, Lindsay T. Johnson, Sean A. Dilliard, and Daniel J. Siegwart. Selective organ targeting (SORT) nanoparticles for tissue-specific mRNA delivery and CRISPR-Cas gene editing. *Nature Nanotechnology*, 15(4):313–320, 2020. doi: 10.1038/s41565-020-0669-6.
- Miruna Cretu, Charles Harris, Ilia Igashov, Arne Schneuing, Marwin Segler, Bruno Correia, Julien Roy, Emmanuel Bengio, and Pietro Liò. SynFlowNet: Design of diverse and novel molecules with synthesis constraints. In *International Conference on Learning Representations*, 2025.
- Pieter R. Cullis and Michael J. Hope. Lipid nanoparticle systems for enabling gene therapies. *Molecular Therapy*, 25(7):1467–1475, 2017. doi: 10.1016/j.ymthe.2017.03.013.
- Wenhao Gao and Connor W. Coley. The synthesizability of molecules proposed by generative models. *Journal of Chemical Information and Modeling*, 60(12):5714–5723, 2020. doi: 10.1021/acs.jcim.0c00174.
- Wenhao Gao, Rocío Mercado, and Connor W. Coley. Amortized tree generation for bottom-up synthesis planning and synthesizable molecular design. In *International Conference on Learning Representations*, 2022.
- Samuel Genheden, Amol Thakkar, Veronika Chadimová, Jean-Louis Reymond, Ola Engkvist, and Esben Bjerrum. AiZynthFinder: A fast, robust and flexible open-source software for retrosynthetic planning. *Journal of Cheminformatics*, 12:70, 2020. doi: 10.1186/s13321-020-00472-1.
- Julian D. Gillmore, Ed Gane, Jorg Taubel, Justin Kao, Marianna Fontana, Michael L. Maitland, Jessica Seitzer, Daniel O’Connell, Kathryn R. Walsh, Kristy Wood, Jonathan

-
- Phillips, Yuanxin Xu, Adam Amaral, Adam P. Boyd, Jeffrey E. Cehelsky, Mark D. McKee, Andrew Schiermeier, Olivier Harari, Andrew Murphy, Christos A. Kyrtasous, Brian Zambrowicz, Randy Soltys, David E. Gutstein, John Leonard, Laura Sepp-Lorenzino, and David Lebwohl. CRISPR-Cas9 in vivo gene editing for transthyretin amyloidosis. *New England Journal of Medicine*, 385(6):493–502, 2021. doi: 10.1056/NEJMoa2107454.
- Xuexiang Han, Hanwen Zhang, Kamila Butowska, Kelsey L. Swingle, Mohamad-Gabriel Alameh, Drew Weissman, and Michael J. Mitchell. An ionizable lipid toolbox for RNA delivery. *Nature Communications*, 12:7233, 2021. doi: 10.1038/s41467-021-27493-0.
- Marine Z. C. Hatit, Curtis N. Dobrowolski, Melissa P. Lokugamage, David Loughrey, Huanzhen Ni, Chiara Zurla, Alejandro J. Da Silva Sanchez, Afsane Radmand, Sebastian G. Huayameres, Ryan Zenhausern, Kalina Paunovska, Hannah E. Peck, Jinwhan Kim, Manaka Sato, Jacob I. Feldman, Michael-Alexander Rivera, Ana Cristian, Yong-Tae Kim, Philip J. Santangelo, and James E. Dahlman. Nanoparticle stereochemistry-dependent endocytic processing improves in vivo mRNA delivery. *Nature Chemistry*, 15(4):508–515, 2023. doi: 10.1038/s41557-023-01138-9.
- Xucheng Hou, Tal Zaks, Robert Langer, and Yizhou Dong. Lipid nanoparticles for mRNA delivery. *Nature Reviews Materials*, 6(12):1078–1094, 2021. doi: 10.1038/s41578-021-00358-0.
- Muthusamy Jayaraman, Steven M. Ansell, Barbara L. Mui, Ying K. Tam, Jianxin Chen, Xinyao Du, David Butler, Laxman Eltepu, Shigeo Matsuda, Jayaprakash K. Narayanan-nair, Kallanthottathil G. Rajeev, Ismail M. Hafez, Akin Akinc, Martin A. Maier, Mark A. Tracy, Pieter R. Cullis, Thomas D. Madden, Muthiah Manoharan, and Michael J. Hope. Maximizing the potency of siRNA lipid nanoparticles for hepatic gene silencing in vivo. *Angewandte Chemie International Edition*, 51(34):8529–8533, 2012. doi: 10.1002/anie.201203263.
- Michał Koziarski, Andrei Rekesh, Dmytro Shevchuk, Almer van der Sloot, Piotr Gaiński, Yoshua Bengio, Cheng-Hao Liu, Mike Tyers, and Robert Batey. RGFN: Synthesizable molecular generation using GFlowNets. In *Advances in Neural Information Processing Systems*, volume 37, pp. 46908–46955. Curran Associates, Inc., 2024. doi: 10.52202/079017-1488.
- Seul Lee, Karsten Kreis, Srimukh Prasad Veccham, Meng Liu, Danny Reidenbach, Yuxing Peng, Saeed Gopal Paliwal, Weili Nie, and Arash Vahdat. GenMol: A drug discovery generalist with discrete diffusion. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 33205–33226. PMLR, 2025.
- Bowen Li, Rajith Singh Manan, Shun-Qing Liang, Akiva Gordon, Allen Jiang, Andrew Varley, Guangping Gao, Robert Langer, Wen Xue, and Daniel Anderson. Combinatorial design of nanoparticles for pulmonary mRNA delivery and genome editing. *Nature Biotechnology*, 41(10):1410–1415, 2023. doi: 10.1038/s41587-023-01679-x.
- Bowen Li, Idris O. Raji, Akiva G. R. Gordon, Lizhuang Sun, Theresa M. Raimondo, Favour A. Oladimeji, Allen Y. Jiang, Andrew Varley, Robert S. Langer, and Daniel G. Anderson. Accelerating ionizable lipid discovery for mRNA delivery using machine learning and combinatorial chemistry. *Nature Materials*, 23(7):1002–1008, 2024. doi: 10.1038/s41563-024-01867-3.
- Kevin T. Love, Kerry P. Mahon, Christopher G. Levins, Kathryn A. Whitehead, William Querbes, J. Robert Dorkin, June Qin, William Cantley, Liu Liang Qin, Timothy Racie, Maria Frank-Kamenetsky, Ka Ning Yip, Rene Alvarez, Dinah W. Y. Sah, Antonin de Fougères, Kevin Fitzgerald, Victor Kotliansky, Akin Akinc, Robert Langer, and Daniel G. Anderson. Lipid-like materials for low-dose, in vivo gene silencing. *Proceedings of the National Academy of Sciences*, 107(5):1864–1869, 2010. doi: 10.1073/pnas.0910603106.

-
- Rohin Maganti, Rahul Maganti, and Mohamad-Gabriel Alameh. Synthesis-constrained discrete diffusion for ionizable lipid generation. In *ICLR 2026 Workshop on Generative and Experimental Perspectives for Biomolecular Design (GEM)*, 2026. URL <https://openreview.net/forum?id=6RFQqfjD06>.
- Yuxuan Ou, Jingyi Zhao, Austin Tripp, Morteza Rasoulboroujeni, and José Miguel Hernández-Lobato. A deep generative model for the design of synthesizable ionizable lipids. In *NeurIPS 2024 Workshop on AI for New Drug Modalities*, 2024. arXiv:2412.00928.
- Fernando P. Polack, Stephen J. Thomas, Nicholas Kitchin, Judith Absalon, Alejandra Gurtman, Stephen Lockhart, John L. Perez, Gonzalo Pérez Marc, Edson D. Moreira, Cristiano Zerbini, Ruth Bailey, Kena A. Swanson, Satrajit Roychoudhury, Kenneth Koury, Ping Li, Warren V. Kalina, David Cooper, Robert W. Frenc, Laura L. Hammitt, Özlem Türeci, Haylene Nell, Axel Schaefer, Serhat Ünal, Dina B. Tresnan, Susan Mather, Philip R. Dormitzer, Uğur Şahin, Kathrin U. Jansen, and William C. Gruber. Safety and efficacy of the BNT162b2 mRNA Covid-19 vaccine. *New England Journal of Medicine*, 383(27): 2603–2615, 2020. doi: 10.1056/NEJMoa2034577.
- Yiming Qin, Manuel Madeira, Dorina Thanou, and Pascal Frossard. DeFoG: Discrete flow matching for graph generation. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 5026–50326. PMLR, 2025.
- Andrei Rekesh, Miruna Cretu, Dmytro Shevchuk, Vignesh Ram Somnath, Pietro Liò, Robert A. Batey, Mike Tyers, Michał Koziarski, and Cheng-Hao Liu. SynCoGen: Synthesizable 3d molecule generation via joint reaction and coordinate modeling. In *International Conference on Learning Representations*, 2026.
- Sean C. Semple, Akin Akinc, Jianxin Chen, Ammen P. Sandhu, Barbara L. Mui, Connie K. Cho, Dinah W. Y. Sah, Derrick Stebbing, Erin J. Crosley, Ed Yaworski, Ismail M. Hafez, J. Robert Dorkin, June Qin, Kieu Lam, Kallanthottathil G. Rajeev, Kim F. Wong, Lloyd B. Jeffs, Lubomir Nechev, Merete L. Eisenhardt, Muthusamy Jayaraman, Mikameh Kazem, Martin A. Maier, Masuna Srinivasulu, Michael J. Weinstein, Qingmin Chen, Rene Alvarez, Scott A. Barros, Soma De, Sandra K. Klimuk, Todd Borland, Verbena Kosovrasti, William L. Cantley, Ying K. Tam, Muthiah Manoharan, Marco A. Ciufolini, Mark A. Tracy, Antonin de Fougères, Ian MacLachlan, Pieter R. Cullis, Thomas D. Madden, and Michael J. Hope. Rational design of cationic lipids for siRNA delivery. *Nature Biotechnology*, 28(2):172–176, 2010. doi: 10.1038/nbt.1602.
- Yue Xu, Shihao Ma, Haotian Cui, Jingan Chen, Shufen Xu, Fanglin Gong, Alex Golubovic, Muye Zhou, Kevin Chang Wang, Andrew Varley, Rick Xing Ze Lu, Bo Wang, and Bowen Li. AGILE platform: A deep learning powered approach to accelerate LNP development for mRNA delivery. *Nature Communications*, 15:6305, 2024. doi: 10.1038/s41467-024-50619-z.
- Lulu Xue, Alex G. Hamilton, Gan Zhao, Zebin Xiao, Rakan El-Mayta, Xuexiang Han, Ningqiang Gong, Xinhong Xiong, Junchao Xu, Christian G. Figueroa-Espada, Sarah J. Shepherd, Alvin J. Mukalel, Mohamad-Gabriel Alameh, Jiaxi Cui, Karin Wang, Andrew E. Vaughan, Drew Weissman, and Michael J. Mitchell. High-throughput barcoding of nanoparticles identifies cationic, degradable lipid-like materials for mRNA delivery to the lungs in female preclinical models. *Nature Communications*, 15(1):1884, 2024. doi: 10.1038/s41467-024-45422-9.

A PRELIMINARIES

This appendix restores the background that the short-paper body compresses.

A.1 IONIZABLE-LIPID STRUCTURE AND MODULAR ASSEMBLY

The structure of an ionizable lipid is typically organized into three functional regions: a protonatable headgroup, a linker, and one or more hydrophobic domains (Semple et al., 2010; Jayaraman et al., 2012). Together, these regions determine the ionization, stability, lipid packing, and membrane interactions of the molecule. Ionizable-lipid libraries are often synthesized by applying a common reaction to different combinations of precursor molecules (Akinc et al., 2008; Love et al., 2010; Li et al., 2024; Xu et al., 2024). Each precursor contributes a defined portion of the final lipid and contains the reactive group required for the assembly. The reaction forms the bonds that connect these precursor-derived regions. As a result, the assembly reaction determines how the regions are connected, while the precursor structures determine the molecular identity of the resulting lipid.

A.2 ASSEMBLY CHEMISTRIES STUDIED

We study three ionizable-lipid assembly chemistries: an amine–aldehyde–isocyanide Ugi three-component reaction, repeated aza-Michael addition, and repeated reductive amination. In the Ugi reaction, an amine head reacts with an aldehyde component and an isocyanide component to form an α -amino amide that connects the headgroup to two hydrophobic regions (Xu et al., 2024). A second assembly strategy uses repeated aza-Michael additions to attach acrylate-bearing hydrophobic components sequentially to reactive amines in the headgroup (Li et al., 2023). Reductive amination sequentially couples aldehyde-bearing hydrophobic components to reactive amines in the headgroup through new carbon–nitrogen bonds (Xue et al., 2024).

A.3 REACTION TRANSFORMS AND EXACT ASSEMBLY

Each assembly chemistry defines a forward transformation that maps role-assigned precursor molecules to the product graphs permitted by that reaction. The corresponding reverse transformation decomposes a complete lipid into the precursor assignments that could have produced it. To verify a decomposition, we computationally replay the reaction on the recovered precursors. The product passes the level-one (L1) assembly check only if the reconstruction matches the original lipid exactly.

A.4 DISCRETE MOLECULAR GRAPH SPACES

Let $\mathcal{A}$ denote a finite vocabulary of atom states and $\mathcal{B}$ a finite vocabulary of bond states. A molecular graph $G = (V, E)$ assigns an atom state $a_i \in \mathcal{A}$ to each node $i \in V$ and a bond state $b_{ij} \in \mathcal{B}$ to each edge $(i, j) \in E$. For a fixed maximum heavy-atom count $N_{\max}$, let $\mathcal{G}$ denote the resulting finite space of labeled molecular graphs.

A.5 DISCRETE FLOW MATCHING

Let $\mathcal{X}$ be a finite state space, p_0 a tractable source distribution, and p_1 the target data distribution. Discrete flow matching constructs a time-indexed probability path $(p_t)_{t \in [0,1]}$ between these distributions, such that $p_{t=0} = p_0$ and $p_{t=1} = p_1$ (Qin et al., 2025). Given a target sample $x_1 \sim p_1$, the linear conditional path is

$$p_t(x_t \mid x_1) = (1 - t)p_0(x_t) + t\delta_{x_1}(x_t), \quad t \in [0, 1], \quad (3)$$

where δ_{x_1} denotes a point mass at x_1 . Marginalizing over x_1 gives

$$p_t(x) = (1 - t)p_0(x) + tp_1(x), \quad (4)$$

so the states remain discrete while their probabilities evolve continuously in time. A continuous-time Markov chain realizes this probability path through transition rates

$R_t(x, y) \geq 0$, where $R_t(x, y)$ is the instantaneous rate of moving from state x to state y . The resulting distribution evolves according to

$$\frac{d}{dt}p_t(x) = \sum_{y \neq x} [p_t(y)R_t(y, x) - p_t(x)R_t(x, y)], \quad (5)$$

which balances the probability flowing into and out of each state. Given an intermediate state x_t , a neural network $p_\theta(x_1 | x_t, t)$ predicts the target state x_1 . We train this predictor by minimizing the denoising cross-entropy

$$\mathcal{L}_{\text{DFM}}(\theta) = \mathbb{E}_{t \sim \mathcal{U}(0,1), x_1 \sim p_1, x_t \sim p_t(\cdot | x_1)} [-\log p_\theta(x_1 | x_t, t)]. \quad (6)$$

For each endpoint x_1 , let $R_t(x, y | x_1)$ realize the conditional path above. We obtain the learned transition rate by averaging these conditional rates under the model prediction:

$$R_t^\theta(x, y) = \sum_{x_1 \in \mathcal{X}} p_\theta(x_1 | x, t) R_t(x, y | x_1). \quad (7)$$

Sampling draws $x_0 \sim p_0$ and evolves the state from $t = 0$ to $t = 1$ according to the learned transition rates R_t^θ .




Figure 3: **FORGE generates complete ionizable-lipid graphs within reaction-program coordinates.** The reaction program specifies how molecular regions participate in an assembly without choosing their structures from a finite inventory. FORGE generates the complete molecular graph, verifies that the specified reaction reconstructs it, and recursively expands unavailable precursors until every branch terminates in evidence-qualified starting materials or the framework abstains. Exact replay and route closure are computational checks and do not establish experimental synthesis success.

B FORGE IN FULL

This appendix gives the complete formulation summarized in Section 2 of the body.

B.1 OVERVIEW

We formulate ionizable-lipid design as conditional generation under a specified reaction program, while leaving precursor identities unspecified. Let $G \in \mathcal{G}$ denote a complete ionizable-lipid graph and $P \in \mathcal{P}$ a reaction program specifying the assembly chemistry, precursor roles, reaction-core positions, and transformation order. FORGE models the conditional distribution $p_\theta(G \mid P)$ without access to precursor identifiers, stored component graphs, or fragment tokens. FORGE is organized into four stages:

1. **Reaction-program specification.** P defines the assembly chemistry, precursor roles, reaction-core positions, transformation order, and coarse per-role morphology, but not precursor identities.
2. **Whole-graph generation.** A conditional discrete flow generates every atom, bond, and topology state of G under P .
3. **Exact assembly verification.** A reverse transform decomposes G into role-assigned precursors, and the corresponding forward transform must reconstruct G exactly.
4. **Recursive synthesis-program construction.** FORGE searches for preparation routes to generated precursors that are not already supported as available starting materials and abstains when a route cannot be completed.

Figure 3 summarizes how FORGE separates reaction semantics from molecular identity and connects whole-graph generation to recursive synthesis-program construction.

B.2 SYNTHESIS-GROUNDED DESIGN TASK

Given a reaction program P , FORGE generates a complete molecular graph G . A downstream bounded program constructor then assesses the exact-assembly precursors and returns a recursive synthesis program T or an explicit abstention outcome. Let $\mathfrak{T}$ denote the space of recursive synthesis programs and $\perp$ an explicit abstention outcome. We write the design process as

$$G \sim p_\theta(\cdot \mid P), \quad T = \mathcal{A}(G, P) \in \mathfrak{T} \cup \{\perp\}. \quad (8)$$

$\mathcal{A}$ denotes the bounded program constructor that verifies the final assembly, searches for preparation routes to unavailable precursors, and checks terminal supplier evidence. P and $\mathcal{A}$ act at different stages: P guides the generation of G , whereas $\mathcal{A}$ assesses the completed graph without modifying it. A pair (G, T) is designated route-complete only when G passes molecular and exact-assembly verification and every required precursor traces to an available starting material; otherwise the route assessor abstains.

B.3 REACTION-PROGRAM SPECIFICATION

Definition B.1 (Reaction program). *A reaction program is a tuple*

$$P = (c, d, \mathbf{r}, \mathbf{k}, M), \quad (9)$$

where c identifies the ordered assembly chemistry, d is the program depth, $\mathbf{r}$ assigns precursor roles, $\mathbf{k}$ marks reaction-core positions, and M specifies coarse per-role morphology. For each precursor role r , M records four integers: the number of exterior atoms n_r , junction budget j_r , cycle rank γ_r , and attachment count a_r .

M sets only coarse structural constraints; the model generates the exact topology, atoms, and bonds of each region. FORGE first samples M from a training-fold morphology prior and then generates G conditioned on the resulting reaction program:

$$p_\theta(G, M | c) = p_{\text{prog}}(M | c) p_\theta(G | M, c). \quad (10)$$

The prior sets each region’s atom count, branching budget, cycle rank, and number of core attachments, but not its molecular structure. A chemistry-specific adapter defines the precursor roles, reaction core, reverse decomposition, and forward transform associated with c . For each graph position i , FORGE combines learned embeddings of the program identity, program depth, precursor role, and reaction-core position:

$$h_i^P = \text{LN}(e_c(c) + e_d(d) + e_r(r_i) + e_k(k_i)). \quad (11)$$

These coordinates tell the generator which precursor role each position plays and whether it lies in the reaction core, but not which molecular structure should occupy that role.

B.4 PROGRAM-DEFINED MOLECULAR SPACE

Reaction programs constrain how a lipid is assembled without specifying the molecular structures that occupy its precursor roles. Let $\mathcal{G}_{\text{bounded}} \subseteq \mathcal{G}$ denote the molecular graphs within the declared atom and bond vocabularies, the 194-heavy-atom limit, and the three-closure limit. We write $G \models M(P)$ when its per-role atom counts, cycle ranks, and core-attachment counts match M , and its junction counts do not exceed the corresponding budgets. The valid, exact-L1 molecular space associated with P is

$$\mathcal{G}_{\text{L1}}(P) = \{G \in \mathcal{G}_{\text{bounded}} : \text{Valid}(G) = 1, G \models M(P), \text{L1}(G, P) = 1\}. \quad (12)$$

A generated lipid can pass the exact L1 check even when one or more recovered precursors never appeared in the training component set, provided the specified reaction reconstructs the lipid from them. For a reaction program P , let $\mathfrak{T}_P$ denote its admissible program inputs, including precursor identities, role assignments, and program depth. The registered forward operator maps each input to a finite set of canonical products:

$$\mathcal{F}_P : \mathfrak{T}_P \longrightarrow 2^{\mathfrak{G}_{\text{N,K}}}. \quad (13)$$

Given role-specific catalogues $\mathcal{C}$, let $\mathfrak{T}_P(\mathcal{C}) \subseteq \mathfrak{T}_P$ contain the admissible program inputs whose components all belong to their corresponding catalogues. The catalogue-reachable product set is

$$\mathcal{G}_{\text{cat}}(P, \mathcal{C}) = \bigcup_{\mathbf{C} \in \mathfrak{T}_P(\mathcal{C})} \mathcal{F}_P(\mathbf{C}). \quad (14)$$

For any graph G , define its complete program-relative preimage as

$$\mathcal{E}_P^*(G) = \{\mathbf{C} \in \mathfrak{T}_P : G \in \mathcal{F}_P(\mathbf{C})\}. \quad (15)$$

Definition B.2 (Strong catalogue escape).

$$\text{Escape}^*(G; P, \mathcal{C}) = \mathbb{K}[\mathcal{E}_P^*(G) \neq \emptyset \wedge \mathcal{E}_P^*(G) \cap \mathfrak{T}_P(\mathcal{C}) = \emptyset]. \quad (16)$$

Theorem E.3 bounds the catalogue-reachable product set while establishing a zero component-novelty ceiling for the catalogue assembler. Strong catalogue escape additionally

requires that no admissible all-catalogue input produces G . The verifier-returned set $\widehat{\mathcal{E}}_P(G)$ defined below is a subset of $\mathcal{E}_P^*(G)$ when every returned trace passes exact forward replay. A novel component in a returned trace establishes strong catalogue escape only when reverse enumeration is exhaustive or direct catalogue reachability is checked independently. FORGE instead parameterizes $p_\theta(G | P)$ over whole graphs in $\mathcal{G}_{\text{bounded}}$, without including any component catalogue $\mathcal{C}_r$ in the generative state. Reaction programs constrain assembly semantics, while the flow determines precursor identity at atom and bond resolution.

B.5 WHOLE-GRAPH GENERATION

FORGE realizes this whole-graph distribution with a discrete flow over atom, bond, and topology states. To preserve the full 194-atom support without an $O(N_{\text{max}}^2)$ edge state, FORGE encodes each graph as a spanning tree plus at most three residual closure edges. Let

$$X_G = (\mathbf{a}, \boldsymbol{\pi}, \mathbf{b}, \mathbf{e}^L, \mathbf{e}^R, \mathbf{z}) \quad (17)$$

denote the categorical graph state, where $\mathbf{a}$ contains atom states, $\boldsymbol{\pi}$ parent pointers, $\mathbf{b}$ spanning-tree bond states, $\mathbf{e}^L$ and $\mathbf{e}^R$ closure-endpoint pointers, and $\mathbf{z}$ closure-bond states. The sparse decoder is surjective onto the declared graph space $\mathfrak{G}_{N_{\text{max}}, K}$: every graph in this space has at least one exact tree-closure encoding (Theorem E.2). For this study, $N_{\text{max}} = 194$ and $K = 3$. This is a statement about architectural support. It does not imply reaction consistency, nonzero probability under a trained model, or synthesis-program completeness.

The chemistry adapter marks any atom, pointer, or bond states fixed by P . FORGE holds these states constant while corrupting the remaining states along the linear probability path and predicting their clean values from (X_t, t, P) . Let $\mathcal{Q} = \{\mathbf{a}, \boldsymbol{\pi}, \mathbf{b}, \mathbf{e}^L, \mathbf{e}^R, \mathbf{z}\}$ denote the six state families optimized during training. Let μ_{bal} denote the frozen training distribution over graph-program pairs, which assigns equal total mass to each reaction program and preserves the source-balanced weights within each program. For each $q \in \mathcal{Q}$, let $m_j^{(q)}(P) \in \{0, 1\}$ indicate whether coordinate j remains variable under P . FORGE minimizes a program-balanced denoising objective over the variable coordinates:

$$\mathcal{L}_{\text{FORGE}}(\theta) = \mathbb{E}_{(G, P) \sim \mu_{\text{bal}}, t \sim \mathcal{U}(0, 1), X_t \sim p_t(\cdot | X_G, P)} \left[\sum_{q \in \mathcal{Q}} \lambda_q \frac{\sum_j m_j^{(q)}(P) \left[-\log p_\theta \left(X_{G,j}^{(q)} | X_t, t, P \right) \right]}{\max \left\{ 1, \sum_j m_j^{(q)}(P) \right\}} \right]. \quad (18)$$

Here, λ_q weights state family q , and a family with no variable coordinates contributes zero loss. The corruption and sampling operators project every intermediate state onto the program-fixed coordinates. This makes those coordinates invariant by construction, while the masked denoising objective learns the conditional clean-state distributions of the remaining coordinates in the population limit (Proposition E.4 and Theorem E.5).

The denoiser is a graph Transformer over the complete sparse state. Each layer applies graph self-attention with relation biases for parent, child and closure edges, then cross-attends to memory tokens representing program identity, program depth, precursor role and reaction-core position. Cross-attention is repeated in every layer so that reaction semantics can influence both local state updates and long-range coordination among precursor-derived regions. Three softly program-routed residual adapters provide family-specific capacity while retaining a shared whole-graph backbone.

The complete training loss augments the flow objective with precursor-role and reaction-core consistency terms:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{FORGE}} + \lambda_{\text{role}} \mathcal{L}_{\text{role}} + \lambda_{\text{core}} \mathcal{L}_{\text{core}}. \quad (19)$$

Present semantic states receive equal mass within each auxiliary term. Training gives every reaction family equal total loss mass and uses deterministic projected conflicting gradient updates (PCGrad) to remove pairwise gradient conflicts without amplifying a nearly converged family’s gradient. Table 5 prespecifies separate interventions for layerwise cross-attention, each auxiliary loss, routed adapters and gradient conflict control.

At inference, FORGE samples X_0 from the source distribution, holds the program-fixed coordinates constant, and simulates $R_t^\theta(\cdot \mid P)$ from $t = 0$ to $t = 1$. FORGE decodes the terminal state X_1 into a molecular graph G and retains it only if it is connected, valence-valid, and passes molecular sanitization.

B.6 EXACT ASSEMBLY VERIFICATION

To ensure that each retained graph satisfies the specified reaction program, FORGE decomposes it into role-assigned precursors and applies the corresponding forward transform.

Definition B.3 (Exact L1 consistency). *For chemistry c , the reverse transform $\mathcal{D}_c$ maps G to a set of role-assigned precursor tuples. Let $\mathcal{D}_P(G) \subseteq \mathcal{D}_c(G)$ contain only those precursor tuples whose role assignments, reaction-core atoms, and morphology satisfy P . The verifier-returned exact set is*

$$\hat{\mathcal{E}}_P(G) = \{\mathbf{C} \in \mathcal{D}_P(G) : G \in \mathcal{F}_P(\mathbf{C})\}. \quad (20)$$

A graph G satisfies exact L1 consistency under P when this set is nonempty:

$$\text{L1}(G, P) = \mathbb{K}[\hat{\mathcal{E}}_P(G) \neq \emptyset]. \quad (21)$$

Post-hoc filtering changes which completed samples are retained but cannot change the proposal distribution’s exact-L1 acceptance probability. Under independent sampling, this probability determines the complete distribution of accepted counts under a fixed attempt budget and the expected cost of obtaining a fixed number of accepted products (Proposition E.6). FORGE reports decomposition coverage separately from forward-replay precision because identifying a candidate disconnection does not guarantee exact reconstruction of G . Decomposition coverage is the fraction of valid graphs for which $\mathcal{D}_c(G) \neq \emptyset$; forward-replay precision is the fraction of returned precursor tuples for which $\mathcal{F}_c(\mathbf{C}) \cong G$. FORGE rejects G when $\mathcal{D}_c(G) = \emptyset$ and marks it as ambiguous when $|\mathcal{D}_c(G)| > 1$.

B.7 RECURSIVE SYNTHESIS-PROGRAM CONSTRUCTION

The final assembly requires access to every recovered precursor. Exact L1 verification identifies these precursors, but it does not establish whether they are available or how to prepare them. For laboratory planning, an exact final assembly is useful only when every required precursor is either available or connected to available starting materials through a supported preparation route. Let $T = (V_M, V_R, E)$ denote a bipartite synthesis tree rooted at G , where each product molecule points to its preparation reaction and each reaction points to its required reactants. FORGE constructs these routes recursively and returns an abstention for route completion when any branch remains unresolved; the exact-L1 generated product itself remains valid. A molecule node terminates when a current, dated supplier record lists an exact constitutional match as available; otherwise, FORGE must expand that node through a supported, forward-verified reaction. FORGE uses a bounded level-two (L2) planner to propose each expansion. For each proposal, FORGE records the reaction, reactants, product, and supporting evidence, and accepts the step only when the forward transform reconstructs the target precursor exactly. Each accepted reaction adds its reactants as new molecule nodes, and FORGE applies the same termination-or-expansion rule recursively. For each L2 reaction node r , let $\text{Verify}(r) = 1$ when its forward transform reconstructs its product exactly. For each leaf molecule ℓ , let $\text{Available}(\ell) = 1$ when it satisfies the terminal supplier-evidence rule. Let $V_R^{\text{L2}}(T)$ denote the L2 reaction nodes and $L(T)$ the leaf molecules in T .

$$\text{Complete}(G, T, P) = \text{L1}(G, P) \prod_{r \in V_R^{\text{L2}}(T)} \text{Verify}(r) \prod_{\ell \in L(T)} \text{Available}(\ell). \quad (22)$$

Let $\mathcal{B}$ fix the route-depth and query budgets, registered transforms, evidence-admission policy, and dated supplier snapshot used by the assessor $\mathcal{A}_\mathcal{B}$. Under finite search bounds and terminating evidence queries, the assessor terminates and returns a dossier only when

Algorithm 1 FORGE generation and fail-closed synthesis-program construction.

Require: trained rates R_t^θ , reaction program P , transition count K , attempt budget B , recorded seed s , and frozen route contract $\mathcal{B}$
Ensure: accepted computational dossiers $\mathcal{S}$ and typed failure ledger $\mathcal{H}$

- 1: initialize the random-number streams from s ; $\mathcal{S} \leftarrow \emptyset$; $\mathcal{H} \leftarrow \emptyset$
- 2: **for** $b = 1, \dots, B$ **do**
- 3: sample $X_0 \sim p_0(\cdot | P)$ and set $X \leftarrow \Pi_P(X_0)$
- 4: **for** $k = 0, \dots, K - 1$ **do**
- 5: $t_k \leftarrow k/K$
- 6: $X \leftarrow \Pi_P(\text{JUMPSTEP}(X, R_{t_k}^\theta(\cdot | X, P)))$
- 7: **end for**
- 8: $G \leftarrow \text{DECODE}(X)$
- 9: **if** G is disconnected, valence-invalid, or fails molecular sanitization **then**
- 10: $\mathcal{H} \leftarrow \mathcal{H} \cup \{(b, \text{INVALIDGRAPH})\}$
- 11: **continue**
- 12: **end if**
- 13: $\mathcal{E}_P(G) \leftarrow \{\mathbf{C} \in \mathcal{D}_P(G) : G \in \mathcal{F}_P(\mathbf{C})\}$
- 14: **if** $\mathcal{E}_P(G) = \emptyset$ **then**
- 15: $\mathcal{H} \leftarrow \mathcal{H} \cup \{(b, \text{L1FAILURE})\}$
- 16: **continue**
- 17: **end if**
- 18: $Y \leftarrow \mathcal{A}_B(G, P)$ $\triangleright$ recursively assess the exact L1 precursors
- 19: **if** Y is a tree T and $\text{Complete}(G, T, P) = 1$ **then**
- 20: $\mathcal{S} \leftarrow \mathcal{S} \cup \{\text{DOSSIER}(G, P, T)\}$
- 21: **else**
- 22: $\mathcal{H} \leftarrow \mathcal{H} \cup \{(b, Y)\}$ $\triangleright$ preserve the typed outcome
- 23: **end if**
- 24: **end for**
- 25: **return** $\mathcal{S}, \mathcal{H}$

$\text{Complete}(G, T, P) = 1$ (Lemma E.7 and Theorem E.8). An abstention means that FORGE did not complete a dossier under $\mathcal{B}$; it does not imply that no synthesis route exists outside the frozen search space or evidence snapshot. The guarantee is computational and does not establish experimental synthesis success.

Algorithm 1 summarizes the complete inference procedure. Every iteration consumes one generation attempt, and FORGE retains the outcome of every failed attempt in a typed ledger.

For every accepted pair (G, T) , FORGE records a computational synthesis dossier containing all molecules and reactions in T , the evidence supporting each reaction, and the dated supplier record for every terminal material. We classify each generated lipid into one of four support tiers according to which elements of its assembly lie outside the fixed reference space: **E0 (enumerated)**: the generated lipid matches a product that was already constructed computationally from known components using a supported reaction. **E1 (known components)**: the generated lipid was not included in the bounded E0 enumeration, but a supported reaction assembles it entirely from known components. **E2 (generated precursor)**: a supported reaction assembles the lipid, but at least one precursor is newly generated and requires a complete L2 route to available starting materials. **E3 (new assembly family)**: the lipid requires a final assembly reaction outside the qualified reaction registry.

C EXTENDED EXPERIMENTAL PROTOCOL AND RESULTS

This appendix gives the full reporting contract, including every matched arm, ablation and diagnostic summarized in Section 3 of the body.

C.1 EXPERIMENTAL SETTINGS

Our evaluation spans the three assembly programs defined above, with Ugi as the deeply supervised case, repeated aza-Michael addition as a second family, and repeated reductive amination as a lighter stress test. The evaluation tests four capabilities: learning multiple assembly programs within one model, using reaction semantics during generation rather than after it, producing exact assemblies beyond a finite component catalogue, and tracing generated precursors to available starting materials. We compare four matched training arms: Ugi-only conditioning, shared conditioning across all three programs, shared training without program information, and shared training with cyclically incorrect program labels. All arms use the same architecture and optimization budget, and we evaluate the prespecified final checkpoint from each of three independent seeds. Each matched comparison uses the same number of generation attempts per program. Our primary metric is exact-L1 yield per generation attempt: the number of generated graphs that pass molecular validity, program-aware decomposition, and exact forward replay divided by the total number of sampling attempts. Secondary metrics separate graph quality (validity and connectedness), reaction recovery (coverage, precision, abstention, and ambiguity), and exploration (unique-product fraction, whole-product and component novelty, molecular diversity, and effective component count). We report means across seeds, individual seed values, and 95% paired-seed bootstrap intervals; generated molecules are not treated as independent replicates. Training uses only the component-disjoint training folds. Calibration data are diagnostic, and held-out products neither enter training nor select checkpoints. The three shared arms assign equal total sampling mass to each program while preserving the frozen within-program source weights. We assess whether shared training preserves Ugi behavior using prespecified non-inferiority margins and require zero program-fixed-state violations and zero support overflows.

Two additional comparisons isolate the role of reaction semantics. First, we compare conditioned generation with null generation followed by exact post-hoc reaction filtering under the same attempt budget; the cyclic-label arm tests whether chemical meaning matters beyond the presence of a categorical input. Second, we compare FORGE with a train-only finite-component assembler that receives the exact training components, supported forward chemistry, source-weighted role marginals, and the same number of attempts. We report reaction validity and exploration separately because the catalogue can achieve high exact-L1 yield while having zero input-component novelty by construction. These comparisons make no L2 planner, biological-oracle, or candidate-selection calls.

For the Ugi program, we additionally define a common external-baseline protocol with one reaction-space comparator, RGFN, and two generic whole-molecule comparators, unconditional DeFoG and GenMol/SAFE (Koziarski et al., 2024; Qin et al., 2025; Lee et al., 2025). The Synthesis-DAG method of Ou et al. (2024), SynFlowNet and SynCoGen remain methodological context rather than empirical rows (Cretu et al., 2025; Reakes et al., 2026). The released Synthesis-DAG repository does not contain the paper-specific ionizable-lipid implementation, the released SynFlowNet action space cannot represent the qualified three-reactant Ugi transform, and the released SynCoGen repository has no usable code license. We do not replace missing functionality with an imitation, alter Ugi into sequential pseudo-reactions, or execute unlicensed code. An external row enters the numerical comparison only after the method has been executed on the same training split, generation-attempt budget, seed set, and common molecular and exact-L1 evaluation pipeline used for FORGE. Published values obtained on the method’s original dataset are not mixed with results from this common evaluation. Methods that do not support the additional assembly programs are compared on Ugi only rather than assigned artificial aza-Michael or reductive-amination results. The matched internal suite separately includes unconditioned generation plus post-hoc filtering, an oracle finite-catalogue assembler, a learned inventory selector, two learned component-factorized generators, and targeted interventions on the Transformer. Table 5 re-

Exact-L1 yield (%) and transform verification across three seeds				
Program	FORGE conditioned	Shared null	Cyclic program	Decomp./replay (%)
Ugi	78.7 $\pm$ 10.1	11.4 $\pm$ 1.1	40.2 $\pm$ 0.7	90.7/100.0
Aza-Michael	77.5 $\pm$ 3.7	4.0 $\pm$ 0.8	3.8 $\pm$ 0.6	100.0/100.0
Reductive amination	31.3 $\pm$ 14.3	13.2 $\pm$ 1.3	18.1 $\pm$ 0.8	84.3/100.0

Shared conditioning improved Ugi exact-L1 yield over Ugi-only training by 53.26 percentage points (95% interval 41.24 to 61.82).

Figure 4: **A shared model learns multiple ionizable-lipid assembly programs.** Results are reported separately for Ugi, repeated aza-Michael addition and repeated reductive amination. Cells show means and sample standard deviations across three independent training seeds. The stated Ugi contrast uses a 95% paired-seed bootstrap interval. The Ugi-only arm is evaluated only on Ugi, and results are not pooled across reaction programs. All values in this figure and Table 3 are hash verified.

Table 3: **Matched production comparison across reaction programs.** Point estimates report the mean across three independent seeds; the exact-L1 column additionally reports the sample standard deviation. Paired contrasts include 95% paired-seed bootstrap intervals; individual seed values appear in the supplement. Validity and connectedness use all generation attempts as the denominator. Decomposition coverage, replay precision, abstention and ambiguity retain their prespecified conditional denominators. Novelty is conditioned on valid products or unique returned decompositions, as appropriate. Fraction pairs and exact-L1 yield are percentages; diversity is mean pairwise ECFP4 distance and N_{eff} is the effective generated-component count. Results are never pooled across reaction programs.

Program	Arm	Valid./Conn.	Decomp./Replay	Abst./Ambig.	Exact-L1/attempt	Whole/Comp. novelty	Diversity/ N_{eff}
Ugi	Ugi-only	28.0/28.0	89.9/100.0	10.1/0.0	25.4 $\pm$ 12.0	99.8/92.4	0.597/134.8
Ugi	FORGE conditioned	86.6/86.6	90.7/100.0	9.3/0.0	78.7 $\pm$ 10.1	99.8/96.2	0.621/320.8
Ugi	Shared null	19.5/19.5	58.8/100.0	41.2/0.0	11.4 $\pm$ 1.1	100.0/99.9	0.754/219.9
Ugi	Cyclic program	43.4/43.4	92.7/100.0	7.3/0.0	40.2 $\pm$ 0.7	99.9/90.1	0.562/138.0
Aza-Michael	FORGE conditioned	77.5/77.5	100.0/100.0	0.0/54.9	77.5 $\pm$ 3.7	99.7/99.2	0.521/156.9
Aza-Michael	Shared null	16.5/16.5	24.3/100.0	75.7/0.1	4.0 $\pm$ 0.8	99.9/99.7	0.635/111.3
Aza-Michael	Cyclic program	25.7/25.7	15.2/100.0	84.8/0.0	3.8 $\pm$ 0.6	99.8/98.3	0.388/51.5
Reductive amination	FORGE conditioned	36.5/36.5	84.3/100.0	15.7/74.1	31.3 $\pm$ 14.3	98.1/100.0	0.624/102.9
Reductive amination	Shared null	15.4/15.4	85.4/100.0	14.6/70.7	13.2 $\pm$ 1.3	99.5/100.0	0.615/77.2
Reductive amination	Cyclic program	21.4/21.4	84.9/100.0	15.1/74.6	18.1 $\pm$ 0.8	96.2/100.0	0.570/34.8

ports the completed three-seed Transformer mechanism and component-factorization study. Each intervention uses the same folds, generation budget and evaluation code as the full Transformer.

C.2 SHARED REACTION-PROGRAM LEARNING ACROSS THREE ASSEMBLY CHEMISTRIES

Across 3,072 held-out generation attempts per program in each of three independent training seeds, the final conditioned model achieved exact-L1 yields of $78.7 \pm 10.1\%$ for Ugi (72.3%, 90.3% and 73.5% by seed), $77.5 \pm 3.7\%$ for repeated aza-Michael addition (76.5%, 81.5% and 74.3%), and $31.3 \pm 14.3\%$ for repeated reductive amination (38.6%, 14.8% and 40.4%) (mean $\pm$ sample standard deviation; Figure 4 and Table 3). Exact forward-replay precision was 100% in every program and seed. However, mean decomposition ambiguity was 54.9% for repeated aza-Michael addition and 74.1% for repeated reductive amination, and the latter program showed substantial between-seed variation. Table 3 reports molecular validity, connectedness, decomposition coverage, forward-replay precision, abstention, ambiguity, novelty and diversity separately for each program.

Relative to the Ugi-only model, shared training changed Ugi exact-L1 yield by 53.26 percentage points (95% paired-seed interval 41.24 to 61.82) and met every prespecified Ugi-retention gate. The final conditioned model produced zero fixed-state violations and zero support overflows across 27,648 held-out attempts (three seeds, three programs and 3,072 attempts per program). Together, these results support the conclusion that one conditional flow can learn the three specified reaction-program coordinates without erasing the deeply supported Ugi behavior.

Table 4: **Method-blind structural-realism diagnostic.** Every entry is mean $\pm$ sample standard deviation across three training seeds and retains all 3,072 requested attempts in the denominator. Support is the declared C/N/O/P/S and 194-heavy-atom space. FP and descriptor P/C report manifold precision per 1,000 attempts and coverage of the source-study-held-out R0 reference in percent. Desc. W is the normalized descriptor Wasserstein distance; lower is closer. C2ST is grouped classifier two-sample area under the receiver operating characteristic curve; 0.5 is indistinguishable. Unique is conditioned on connected outputs. N_{eff} is the Shannon effective molecule count. N/E denotes a metric that cannot be estimated. This post-hoc structural diagnostic is not a property, activity or synthesis-success

Method	Connected/1k	Support/1k	FP P/C	Desc. P/C	Desc. W	C2ST	Unique (%)	N_{eff}	Diversity
RGFN	0.0 $\pm$ 0.0	0.0 $\pm$ 0.0	0.0 $\pm$ 0.0/0.00 $\pm$ 0.00	0.0 $\pm$ 0.0/0.00 $\pm$ 0.00	N/E	N/E	N/E	0.0 $\pm$ 0.0	N/E
DeFoG unconditional	346.5 $\pm$ 42.3	334.2 $\pm$ 39.1	0.0 $\pm$ 0.0/0.00 $\pm$ 0.00	17.3 $\pm$ 1.4/1.45 $\pm$ 0.06	2.067 $\pm$ 0.009	0.961 $\pm$ 0.007	100.0 $\pm$ 0.0	1064.3 $\pm$ 129.9	0.726 $\pm$ 0.013
GenMol/SAFE	627.5 $\pm$ 6.2	627.4 $\pm$ 6.2	0.0 $\pm$ 0.0/0.00 $\pm$ 0.00	0.0 $\pm$ 0.0/0.00 $\pm$ 0.00	3.199 $\pm$ 0.003	N/E	0.9 $\pm$ 0.1	3.2 $\pm$ 0.1	0.873 $\pm$ 0.009
Finite catalogue oracle	1000.0 $\pm$ 0.0	956.5 $\pm$ 2.2	0.2 $\pm$ 0.2/0.05 $\pm$ 0.05	54.6 $\pm$ 4.4/2.02 $\pm$ 0.17	2.107 $\pm$ 0.002	0.971 $\pm$ 0.007	83.8 $\pm$ 0.8	2395.5 $\pm$ 31.7	0.720 $\pm$ 0.000
Learned inventory selector	1000.0 $\pm$ 0.0	945.0 $\pm$ 6.5	0.1 $\pm$ 0.2/0.02 $\pm$ 0.03	56.6 $\pm$ 7.2/1.97 $\pm$ 0.46	2.101 $\pm$ 0.005	0.974 $\pm$ 0.002	85.0 $\pm$ 1.4	2443.0 $\pm$ 55.8	0.718 $\pm$ 0.001
FORGE Transformer	866.0 $\pm$ 90.2	864.6 $\pm$ 91.2	0.0 $\pm$ 0.0/0.00 $\pm$ 0.00	8.8 $\pm$ 2.3/1.16 $\pm$ 0.20	2.089 $\pm$ 0.005	0.973 $\pm$ 0.019	95.7 $\pm$ 1.9	2479.4 $\pm$ 177.2	0.621 $\pm$ 0.013

C.3 REACTION SEMANTICS ALTER GENERATION RATHER THAN MERELY FILTERING ITS OUTPUT

Under matched attempt budgets, conditioned generation produced 786.8, 774.6 and 312.7 exact-L1 products per 1,000 attempts for Ugi, repeated aza-Michael and repeated reductive amination, compared with 114.4, 39.7 and 131.7 for matched null generation followed by post-hoc filtering (Figure 5a–c). The paired conditioned-minus-post-hoc differences were 67.24 percentage points for Ugi (95% interval 60.84 to 79.95), 73.49 for aza-Michael (69.69 to 78.45), and 18.10 for reductive amination (2.51 to 26.04).

The cyclic control retained precursor roles, reaction-core coordinates and program depth but assigned the wrong reaction identity. Relative to correct conditioning, this intervention changed exact-L1 yield by 38.43 percentage points for Ugi (95% interval 32.00 to 49.41), 73.67 for aza-Michael (71.13 to 77.02), and 13.15 for reductive amination (-3.91 to 22.01). Forward-replay precision remained 100% for every returned decomposition in these arms.

A separate inference-time intervention held the trained network and Ugi request fixed while changing only program coordinates. The factual exact-L1 yield was 39.69%. Nulling all program coordinates reduced it to 3.62%, a 36.07-point loss (95% interval 34.28 to 37.24). Reversing the role assignment caused a 4.49-point loss (1.37 to 8.27), and substituting the reductive-amination program token caused a 22.60-point loss (14.06 to 27.47). The aza-Michael-token contrast was unresolved at 2.00 points (-2.18 to 5.60). Thus the coordinates affect generation, but the evidence does not support an equal or individually identifiable effect for every coordinate.

C.4 WHOLE-GRAPH GENERATION PRODUCES PRECURSOR IDENTITIES BEYOND A FINITE VOCABULARY

Table 2 defines the common Ugi benchmark used to compare FORGE with external synthesis-aware generators, matched internal controls, and the finite-catalogue oracle. Every yield is normalized by the number of generation attempts, including attempts that fail before producing a valid molecule. All included methods were run under the common split, attempt budget, seed set and verifier.

The learned finite-inventory selector achieved 1000.0 $\pm$ 0.0 exact-L1 products per 1,000 attempts, but its finite action space imposed zero open-ended component generation. FORGE produced 786.8 $\pm$ 100.8 exact-L1 products and 751.2 $\pm$ 80.2 unique method-visible open-ended exact-L1 products per 1,000 attempts. Among open-ended external generators, unconditional DeFoG produced 280.6 $\pm$ 26.6 exact-L1 products per 1,000 attempts and recovered held components at 15.5 $\pm$ 2.4 per 1,000 attempts, compared with 0.1 $\pm$ 0.2 for FORGE. GenMol/SAFE produced 627.5 $\pm$ 6.2 valid connected outputs but no exact Ugi-L1 product, while RGFN emitted no admissible molecular output under the strict native contract. These outcomes expose separate tradeoffs between reaction consistency, open-ended identity generation and recovery of specifically reserved components.

Table 5: **Transformer architecture and mechanism ablations.** The five Transformer intervention rows change only the named mechanism from the full model; the two FACT rows instead generate precursor roles independently at matched or increased capacity. All arms use identical data folds, corruption, context schedule, checkpoint rule and sampling budget. Entries report the mean and [95% seed-bootstrap interval] across three independent training seeds. Held-out loss is the mean denoising loss at $t = 0.5$ across the three component-disjoint program partitions. Cross-role fidelity is the Frobenius distance to the held-out residualized descriptor matrix, for which lower values indicate closer agreement. Exact-L1 yields use all 3,072 attempts per program and seed as the denominator.

Arm	Held-out loss	Ugi L1/1k	BL L1/1k	LX L1/1k	Held-comp. L1/1k	Cross-role fidelity	Diversity
Input-only program	8.48 [7.72, 9.14]	737.30 [631.18, 845.70]	779.84 [757.49, 820.31]	277.34 [230.14, 339.52]	0.11 [0.00, 0.33]	1.79 [1.70, 1.91]	0.60 [0.59, 0.61]
No role loss	7.87 [7.78, 8.00]	771.59 [672.20, 889.00]	772.14 [743.49, 820.64]	289.28 [115.23, 396.16]	0.11 [0.00, 0.33]	1.94 [1.89, 1.98]	0.59 [0.58, 0.60]
No core loss	7.85 [7.78, 7.98]	753.69 [674.15, 906.58]	772.79 [726.89, 819.34]	300.46 [151.69, 409.18]	0.11 [0.00, 0.33]	1.89 [1.75, 1.98]	0.59 [0.58, 0.59]
No routed adapters	7.81 [7.44, 8.04]	718.64 [592.45, 875.33]	779.08 [744.47, 823.89]	326.71 [265.62, 384.11]	0.00 [0.00, 0.00]	1.88 [1.75, 2.06]	0.59 [0.59, 0.59]
No gradient conflict control	7.98 [7.82, 8.20]	798.07 [732.42, 916.02]	768.88 [721.68, 806.32]	326.93 [149.09, 427.73]	0.11 [0.00, 0.33]	1.96 [1.87, 2.09]	0.59 [0.58, 0.59]
FACT-matched	7.09 [6.86, 7.34]	51.97 [37.76, 68.68]	393.66 [327.47, 460.94]	509.87 [457.36, 562.83]	0.00 [0.00, 0.00]	1.45 [1.14, 1.72]	0.56 [0.55, 0.57]
FACT-generous	8.24 [8.03, 8.60]	296.66 [161.78, 508.79]	686.96 [584.64, 811.52]	561.31 [491.54, 635.74]	0.11 [0.00, 0.33]	1.16 [1.14, 1.18]	0.60 [0.60, 0.60]
FORGE Transformer	7.88 [7.68, 8.12]	791.78 [725.91, 916.67]	762.48 [725.26, 815.10]	318.14 [134.44, 413.41]	0.11 [0.00, 0.33]	1.94 [1.76, 2.09]	0.59 [0.58, 0.59]

No evaluated method reproduced the full held-out observed-lipid distribution: where estimable, grouped C2ST AUC remained between 0.961 and 0.974, and fingerprint-manifold precision was zero or near zero (Table 4). GenMol/SAFE’s high pairwise diversity was paired with an effective count of only 3.2 ± 0.1 molecules, confirming severe frequency collapse. FORGE retained 2479.4 ± 177.2 effective molecules, but only 8.8 ± 2.3 attempts per 1,000 entered the descriptor manifold. We therefore interpret this analysis as a multi-axis structural diagnostic rather than a single measure of whether a generated lipid is reasonable.

C.5 JOINT WHOLE-LIPID DENOISING HAS REACTION-DEPENDENT ADVANTAGES OVER ROLE ISOLATION

The full Transformer produced 791.8 Ugi and 762.5 repeated aza-Michael exact-L1 products per 1,000 attempts, compared with 52.0 and 393.7 for the parameter-matched role-isolated FACT model (Table 5). The full Transformer and FACT-matched model each had 5.33 million trainable parameters; FACT-generous had 29.40 million. The higher-capacity FACT control improved to 296.7 for Ugi and 687.0 for repeated aza-Michael addition, but remained below the full Transformer in every seed for both programs. The direction reversed for repeated reductive amination: the full Transformer produced 318.1 exact-L1 products per 1,000 attempts, whereas FACT-matched and FACT-generous produced 509.9 and 561.3. Thus, joint whole-lipid denoising improved exact assembly for Ugi and repeated aza-Michael addition but was not uniformly superior across reaction programs. Because the prespecified conditional-dependence diagnostic was negative, this architecture contrast does not establish detectable cross-role dependence in the present data.

Removing layerwise program cross-attention reduced mean exact-L1 yield by 54.5 products per 1,000 attempts for Ugi and 40.8 for repeated reductive amination, while increasing repeated aza-Michael yield by 17.4. Removing the role or core objective and removing routed adapters also produced program-dependent changes rather than a uniform loss. The arm without gradient-conflict control slightly exceeded the full Transformer mean in all three programs. We therefore do not claim that every introduced module independently improves exact-L1 yield. These one-at-a-time interventions identify which effects replicate across chemistries; they are not pooled into a composite architecture score.

At the matched sampling budget, FORGE produced 751.2 unique method-visible open-ended exact-L1 Ugi products per 1,000 attempts, averaged across the three independent seeds. The corresponding values were 313.9 for aza-Michael addition and 34.8 for reductive amination. The finite-component assembler produced zero by construction because every sampled product uses only components from the training catalogue (Figure 5d–f and Table 6). The catalogue achieved exact-L1 yields of 1,000.0, 723.2 and 1,000.0 products per 1,000 attempts, compared with 786.8, 774.6 and 312.7 for FORGE. The corresponding raw validities were 1,000.0, 723.2 and 1,000.0 for the catalogue and 866.0, 774.6 and 365.0 for FORGE. We report these quantities separately because reaction validity and verifier-relative component novelty measure different capabilities.

In a distinct-product Ugi role audit, 19,520 of 26,235 products (74.4%) contained at least one verifier-returned component constitution absent from the training registry. Across all Ugi products admitted to the novelty audit, 21,960 of 26,235 distinct products (83.7%) were absent from the full 112,386-product constitutional reference corpus. FORGE effective component counts were 320.8, 156.9 and 102.9, compared with 86.4, 23.4 and 16.9 for the finite catalogue. The Ugi catalogue contained 185 amines, 75 aldehydes and 35 isocyanides and covered 100% of training products but 0% of the calibration and component-disjoint held-out products. Its Cartesian tuple ceilings were 485,625 for Ugi, 774 for aza-Michael and 490 for reductive amination; these count possible component-program inputs rather than unique forward products. These results establish the intended tradeoff: the exact finite catalogue has high transform consistency, whereas whole-graph generation can produce exact-L1 products whose recovered precursor identities are not restricted to that inventory.

Exact assembly versus the finite-catalogue ceiling				
Program	FORGE L1/1k	Catalogue L1/1k	FORGE open-ended/1k	Catalogue open-ended/1k
Ugi	786.8	1000.0	729.3	0.0
Aza-Michael	774.6	723.2	313.9	0.0
Reductive amination	312.7	1000.0	34.8	0.0

Nulling all program coordinates reduced factual Ugi exact-L1 yield from 39.69% to 3.62%; substituting the reductive-amination token reduced it to 17.09%.

Figure 5: **Reaction conditioning improves sampling efficiency without imposing a finite component vocabulary.** The inset reports inference-time coordinate interventions, and the table compares FORGE with a train-only finite-component assembler that receives exact forward chemistry and the same per-program attempt budget. Values are means across three independent training seeds. The comparisons use identical molecular support and final reaction adapters and make no L2-planner, biological-oracle or candidate-selection calls. The finite catalogue has zero open-ended yield by construction.

Table 6: **Whole-graph generation versus a finite-component assembler.** Exact-L1, unique exact-L1 and open-ended exact-L1 yields are reported per 1,000 generation attempts. An operationally open-ended product has exactly one verifier-returned exact-L1 decomposition and contains at least one component constitution absent from the training catalogue. This verifier-relative metric is not called strong catalogue escape unless reverse enumeration is shown to be exhaustive. The catalogue has zero input-component novelty by construction, but may achieve higher raw validity or exact-L1 yield. FORGE-minus-catalogue contrasts use the three seeds as paired resampling units.

Program	Method	Valid/1k	Exact L1/1k	Unique L1/1k	Open-ended/1k	Whole nov.	Comp. nov.	Diversity/ N_{eff}
Ugi	FORGE	866.0	786.8	751.2	729.3	99.8	96.2	0.621/320.8
Ugi	Finite catalogue	1000.0	1000.0	837.6	0.0	16.9	0.0	0.719/86.4
Aza-Michael	FORGE	774.6	774.6	738.0	313.9	99.7	99.2	0.521/156.9
Aza-Michael	Finite catalogue	723.2	723.2	147.5	0.0	39.8	0.0	0.402/23.4
Reductive amination	FORGE	365.0	312.7	279.7	34.8	98.1	100.0	0.624/102.9
Reductive amination	Finite catalogue	1000.0	1000.0	52.1	0.0	44.8	0.0	0.514/16.9

C.6 NATIVE RECOVERY OF HELD COMPONENTS REMAINS WEAK

The component-disjoint Ugi fold was never used for training, checkpoint selection or inventory construction. Under the common native-generation assessment, the full Transformer recovered only 1 exact-L1 product containing a designated held component in 9,216 attempts, or 0.11 products per 1,000 attempts averaged across three seeds. This is a negative generalization result. It does not conflict with the broader verifier-relative novelty measurements above: generating a component absent from the training registry is easier than recovering one of the specifically reserved held-component identities. We therefore do not claim strong held-component recovery, role-specific held-component coverage, or zero-shot held-reaction-family generalization.

Table 7: **Predicted HeLa messenger-transfection potential under matched generation budgets.** The support-enriched arm uses the already promoted applicability proposal; the nested arm reallocates morphology mass using the frozen weak potency signal. Conservative-high counts are computational predictions, not transfection measurements. The comparison uses the same generator-call and oracle budgets.

Arm	Attempts	Exact L1	Eligible before budget	Oracle budget	Unique conservative-high
Support-enriched	3,072	2955	90	38	27
Nested potency proposal	3,072	2962	82	38	23

C.7 RECURSIVE SYNTHESIS PROGRAMS CONNECT GENERATED LIPIDS TO EVIDENCE-QUALIFIED MATERIALS

FORGE includes a bounded post-generation route assessor for generated precursors that are not already supported as available materials. The assessor first queries the frozen route registry and may invoke single-step or multistep proposal engines where the registry is silent. A proposed transformation is accepted only when the corresponding forward reaction reconstructs the intended precursor under the declared verification contract, and a complete dossier is returned only when every terminal branch closes under the frozen material-evidence policy. Otherwise the assessor abstains. Abstention therefore records insufficient evidence under the declared search and evidence contract, not proof that the molecule is unsynthesizable.

We do not report a separate AiZynthFinder row: a proposal-only route is a hypothesis and cannot change exact evidence closure until every step passes the same forward-verification and terminal-evidence gates (Genheden et al., 2020).

An L2 step is admitted only when its evidence tier permits the claimed reaction scope and its forward transform reconstructs the intended precursor. A terminal branch closes only with a current, dated supplier record for an exact constitutional match. Computational route closure establishes transform consistency and evidence-qualified terminal materials, not experimental synthesis success.

C.8 GUIDANCE DIAGNOSTICS DEFINE THE LIMITS OF ADDITIONAL OPTIMIZATION

At synthesis-guidance strength 0.25, the matched controller evaluated 64 terminal particles and retained 56 canonical representatives per arm. It changed only 8 terminal identities and returned 2 unique route-ready products, compared with 3 after post-hoc assessment. Only five of 48 four-particle morphology groups contained mixed route-readiness values, leaving most ancestry decisions uniform. This negative result does not support promoting mid-trajectory synthesis tilting. The production method therefore remains complete generation followed by proposal-augmented L1/L2/L3 route assessment.

We evaluate biological guidance only in the bounded HeLa messenger-transfection-potential (mTP) diagnostic. A group-disjoint preterminal classifier detected weak morphology-level signal, with ROC AUC 0.623 (95% interval 0.565 to 0.679) and a 1.177-fold observed top-quartile gain. We then compared a frozen applicability-enriched morphology proposal with a nested potency-enriched proposal using 3,072 generator calls per arm. Table 7 reports the matched terminal outcome.

The nested-minus-support difference was -1.30 conservative-high products per 1,000 generator calls (95% interval -2.28 to 0.98). The promotion gate failed. These computations evaluate predictor behavior within a declared applicability domain; they do not establish improved biological activity, delivery or experimental potency.

D EXTENDED COMPUTATIONAL REPORTING TABLES

The following tables define the complete reporting contract for the common Ugi benchmark. Completed matched controls, the learned inventory selector and three external generators are populated from hash-pinned three-seed results.

Table 8: **Training and sampling parity for the common Ugi benchmark.** Generator calls and attempted products are reported separately from reaction-predictor, verifier, planner, and oracle calls. A method is not described as compute-matched when any required quantity is unavailable. N/R means not recorded by the frozen run; N/P means not applicable.

Method	Params.	Train steps	GPU-hours	Sample steps	Attempts	Reaction calls	Route calls	Oracle calls
RGFN	N/R	5,002 iter.	N/R	N/R	$3 \times 3,072$	0	0	0
DeFoG unconditional	N/R	50,000	N/R	1,000	$3 \times 3,072$	0	0	0
GenMol/SAFE	N/R	50,000	N/R	N/R	$3 \times 3,072$	0	0	0
Learned inventory selector	170,855	2,000	N/R	N/P	$3 \times 3,072$	$3 \times 3,072$	0	0
Finite catalogue oracle	N/P	N/P	N/R	N/P	$3 \times 3,072$	N/R	0	0
Shared-null + post-hoc	N/R	N/R	N/R	N/R	$3 \times 3,072$	N/R	0	0
Input-only program	5,328,301	1,700	N/R	N/R	$3 \times 3,072$	N/R	0	0
No role loss	5,328,301	1,700	N/R	N/R	$3 \times 3,072$	N/R	0	0
No core loss	5,328,301	1,700	N/R	N/R	$3 \times 3,072$	N/R	0	0
No routed adapters	5,026,843	1,700	N/R	N/R	$3 \times 3,072$	N/R	0	0
No gradient conflict control	5,328,301	1,700	N/R	N/R	$3 \times 3,072$	N/R	0	0
FACT-matched	5,328,301	1,700	N/R	N/R	$3 \times 3,072$	N/R	0	0
FACT-generous	29,395,993	1,700	N/R	N/R	$3 \times 3,072$	N/R	0	0
FORGE Transformer	5,328,301	1,700	N/R	N/R	$3 \times 3,072$	N/R	0	0

Table 9: **Seed-level common Ugi benchmark results.** Rows are expanded once a method has a hash-verified three-seed result vector. The main table reports the corresponding mean $\pm$ sample standard deviation; generated molecules are not treated as independent training replicates.

Method	Seed(s)	Valid/1k	Exact L1/1k	Unique L1/1k	Open-ended/1k	Held-comp. L1/1k	Diversity
RGFN	20260825	0.0	0.0	0.0	0.0	0.0	N/E
RGFN	20260826	0.0	0.0	0.0	0.0	0.0	N/E
RGFN	20260827	0.0	0.0	0.0	0.0	0.0	N/E
DeFoG unconditional	20260825	299.2	251.0	251.0	251.0	13.7	0.711
DeFoG unconditional	20260826	359.7	288.4	288.4	288.4	18.2	0.733
DeFoG unconditional	20260827	380.5	302.4	302.4	302.4	14.6	0.736
GenMol/SAFE	20260825	627.3	0.0	0.0	0.0	0.0	0.883
GenMol/SAFE	20260826	633.8	0.0	0.0	0.0	0.0	0.869
GenMol/SAFE	20260827	621.4	0.0	0.0	0.0	0.0	0.866
Learned inventory selector	20260825	1000.0	1000.0	863.0	0.0	0.0	0.717
Learned inventory selector	20260826	1000.0	1000.0	835.6	0.0	0.0	0.717
Learned inventory selector	20260827	1000.0	1000.0	849.9	0.0	0.0	0.715
Finite catalogue oracle	20260825	1000.0	1000.0	841.5	0.0	0.0	0.719
Finite catalogue oracle	20260826	1000.0	1000.0	843.1	0.0	0.0	0.719
Finite catalogue oracle	20260827	1000.0	1000.0	828.1	0.0	0.0	0.719
Shared-null + post-hoc	20260825	183.9	114.3	109.7	109.7	0.3	0.738
Shared-null + post-hoc	20260826	187.5	103.5	98.3	98.3	0.7	0.763
Shared-null + post-hoc	20260827	212.6	125.3	117.5	117.5	0.3	0.761
FACT-matched	20260825	59.2	49.5	48.8	48.8	0.0	0.676
FACT-matched	20260826	81.7	68.7	66.1	66.1	0.0	0.658
FACT-matched	20260827	41.3	37.8	37.4	37.4	0.0	0.668
FACT-generous	20260825	391.0	161.8	161.1	161.1	0.3	0.782
FACT-generous	20260826	747.1	508.8	488.3	488.3	0.0	0.768
FACT-generous	20260827	363.9	219.4	215.5	215.5	0.0	0.771
FORGE Transformer	20260825	801.4	722.7	702.1	702.1	0.0	0.617
FORGE Transformer	20260826	969.1	903.0	843.8	843.8	0.3	0.613
FORGE Transformer	20260827	827.5	734.7	707.7	707.7	0.0	0.634

Table 10: **Reaction decomposition outcomes under the common Ugi verifier.** Coverage is conditioned on molecular validity; replay precision is conditioned on returned decompositions. Abstention and ambiguity retain their prespecified denominators. N/E denotes a conditional metric with no eligible valid product or returned decomposition.

Method	Valid products	Decomp. coverage	Replay precision	Abstention	Ambiguity
RGFN	0.0 $\pm$ 0.0	N/E	N/E	N/E	N/E
DeFoG unconditional	1064.3 $\pm$ 129.9	81.2 $\pm$ 2.4	100.0 $\pm$ 0.0	18.8 $\pm$ 2.4	0.0 $\pm$ 0.0
GenMol/SAFE	1927.7 $\pm$ 19.0	0.0 $\pm$ 0.0	N/E	100.0 $\pm$ 0.0	0.0 $\pm$ 0.0
Learned inventory selector	3072.0 $\pm$ 0.0	100.0 $\pm$ 0.0	100.0 $\pm$ 0.0	0.0 $\pm$ 0.0	0.0 $\pm$ 0.0
Finite catalogue oracle	3072.0 $\pm$ 0.0	100.0 $\pm$ 0.0	100.0 $\pm$ 0.0	0.0 $\pm$ 0.0	0.0 $\pm$ 0.0
Shared-null + post-hoc	598.0 $\pm$ 47.9	58.8 $\pm$ 3.5	100.0 $\pm$ 0.0	41.2 $\pm$ 3.5	0.0 $\pm$ 0.0
FACT-matched	186.7 $\pm$ 62.1	86.3 $\pm$ 4.4	100.0 $\pm$ 0.0	13.7 $\pm$ 4.4	0.0 $\pm$ 0.0
FACT-generous	1538.0 $\pm$ 656.9	56.6 $\pm$ 13.7	100.0 $\pm$ 0.0	43.4 $\pm$ 13.7	0.0 $\pm$ 0.0
FORGE Transformer	2660.3 $\pm$ 277.1	90.7 $\pm$ 2.2	100.0 $\pm$ 0.0	9.3 $\pm$ 2.2	0.0 $\pm$ 0.0

Table 11: **Finite component inventories and source-product coverage.** Inventory size counts unique constitutional components visible to the method at generation time. Product coverage requires that every required role component is present in that inventory. The held-component fold is never used to construct an inventory.

Method	Amine	Aldehyde	Isocyanide	Train coverage	Calibration coverage	Held-comp. coverage
RGFN	185	75	35	100.0	0.0	0.0
DeFoG unconditional	N/P	N/P	N/P	N/P	N/P	N/P
GenMol/SAFE	N/P	N/P	N/P	N/P	N/P	N/P
Finite catalogue oracle	185	75	35	100.0	0.0	0.0
Learned inventory selector	185	75	35	100.0	0.0	0.0
Shared-null + post-hoc	N/P	N/P	N/P	N/P	N/P	N/P
FACT-matched	N/P	N/P	N/P	N/P	N/P	N/P
FACT-generous	N/P	N/P	N/P	N/P	N/P	N/P
FORGE Transformer	N/P	N/P	N/P	N/P	N/P	N/P

E THEORETICAL PROPERTIES

This appendix establishes five properties of FORGE. We show that the sparse state space represents every connected labeled graph within the declared vocabulary, size, and cycle-rank support; characterize the reachability ceiling of a finite-component assembler; prove invariance and population consistency for program-masked denoising; quantify the efficiency of post-hoc rejection under a fixed sampling budget; and establish termination and soundness for the bounded recursive route assessor.

These results concern representation, population objectives, and computational verification. They do not imply that finite-data training assigns positive probability to every representable graph, that registered transformations cover arbitrary chemistry, that bounded route search is complete, or that a returned dossier guarantees experimental synthesis success.

E.1 NOTATION AND SETUP

All graphs in this appendix are finite, connected, undirected, loop-free labeled graphs. Let Σ_V and Σ_E be the finite atom-state and bond-state vocabularies, and define

$$\mathfrak{G}_{N,K} = \{G = (V, E) : |V| \leq N, \beta(G) = |E| - |V| + 1 \leq K, \ell_V(V) \subseteq \Sigma_V, \ell_E(E) \subseteq \Sigma_E\}. \quad (23)$$

Graphs are identified up to label-preserving isomorphism.

A sparse FORGE state with n active nodes is well formed when:

1. $1 \leq n \leq N$, and every active node has a label in Σ_V ;
2. the root has no parent and every node $i > 1$ has one parent $\pi_i < i$;
3. every parent edge has a label in Σ_E ;
4. it contains $k \leq K$ distinct closure edges, each carrying a label in Σ_E , joining two active nodes, and duplicating neither a tree edge nor another closure edge; and
5. closure edges are ordered lexicographically by their endpoint pairs.

Let $\mathfrak{X}_{N,K}$ denote this state space and $\text{Dec} : \mathfrak{X}_{N,K} \rightarrow \mathfrak{G}_{N,K}$ its decoder.

E.2 SPARSE GRAPH REPRESENTATION

Lemma E.1 (Tree residual count). *Let $G = (V, E)$ be a connected finite graph and let T be any spanning tree of G . Then*

$$|E \setminus E(T)| = |E| - |V| + 1 = \beta(G). \quad (24)$$

Proof. A spanning tree on $|V|$ vertices contains exactly $|V| - 1$ edges. Because $E(T) \subseteq E$,

$$|E \setminus E(T)| = |E| - |E(T)| = |E| - (|V| - 1) = |E| - |V| + 1. \quad (25)$$

The final expression is the cyclomatic number of a connected graph. $\square$

Theorem E.2 (Bounded representational completeness). *The decoder*

$$\text{Dec} : \mathfrak{X}_{N,K} \longrightarrow \mathfrak{G}_{N,K} \quad (26)$$

is surjective up to label-preserving graph isomorphism. Equivalently:

1. for every $G \in \mathfrak{G}_{N,K}$, there exists $X \in \mathfrak{X}_{N,K}$ such that $\text{Dec}(X) \cong G$; and
2. every $X \in \mathfrak{X}_{N,K}$ decodes to a graph in $\mathfrak{G}_{N,K}$.

Proof. Let $G = (V, E) \in \mathfrak{G}_{N,K}$, and choose a spanning tree $T \subseteq G$. Choose a root of T and enumerate its vertices $v_1, \dots, v_n$ in any root-first traversal. Every non-root vertex then appears after its parent.

Assign node i the label of v_i . For each $i > 1$, let $\pi_i < i$ be the index of the parent of v_i in T , and assign the corresponding tree-edge label. By Lemma E.1, the residual edge set $E \setminus E(T)$ contains $\beta(G) \leq K$ edges. Orient each residual edge by increasing endpoint index, attach its original edge label, and sort the resulting triples lexicographically by endpoint pair. The resulting state is an element of $\mathfrak{X}_{N,K}$.

The decoder reconstructs every edge of T from the parent pointers and every edge of $E \setminus E(T)$ from the closure triples. It also preserves all node and edge labels. The map $i \mapsto v_i$ is therefore a label-preserving isomorphism from $\text{Dec}(X)$ to G .

Conversely, let $X \in \mathfrak{X}_{N,K}$ have n active nodes. Since $\pi_i < i$, repeatedly following parent pointers from any node strictly decreases its index and must terminate at the root. The $n - 1$ parent edges therefore form a connected acyclic graph, hence a spanning tree. If the state contains $k \leq K$ closure edges, the decoded graph has

$$|E| = (n - 1) + k \quad (27)$$

distinct edges. Its cyclomatic number is consequently

$$|E| - |V| + 1 = (n - 1 + k) - n + 1 = k \leq K. \quad (28)$$

The remaining well-formedness conditions preserve the declared node and edge vocabularies and exclude loops or duplicate edges. Thus $\text{Dec}(X) \in \mathfrak{G}_{N,K}$. $\square$

E.3 FINITE-COMPONENT REACHABILITY

Theorem E.3 (Finite-component reachability). *Assume that $\mathfrak{T}_P(\mathcal{C})$ is finite and that every forward-outcome set is finite. Then*

$$|\mathcal{G}_{\text{cat}}(P, \mathcal{C})| \leq \sum_{\mathbf{C} \in \mathfrak{T}_P(\mathcal{C})} |\mathcal{F}_P(\mathbf{C})|. \quad (29)$$

If $\mathfrak{T}_P(\mathcal{C}) \neq \emptyset$ and

$$M_P(\mathcal{C}) = \max_{\mathbf{C} \in \mathfrak{T}_P(\mathcal{C})} |\mathcal{F}_P(\mathbf{C})|, \quad (30)$$

then

$$|\mathcal{G}_{\text{cat}}(P, \mathcal{C})| \leq M_P(\mathcal{C}) |\mathfrak{T}_P(\mathcal{C})|. \quad (31)$$

Moreover,

$$G \in \mathcal{G}_{\text{cat}}(P, \mathcal{C}) \iff \mathcal{E}_P^*(G) \cap \mathfrak{T}_P(\mathcal{C}) \neq \emptyset. \quad (32)$$

Every catalogue attempt uses only catalogue components, so its attainable component-novelty fraction is exactly zero.

Proof. The reachable set is a finite union of forward-outcome sets. The cardinality of a union is at most the sum of the cardinalities of its members, which proves the first inequality. Bounding each summand by $M_P(\mathcal{C})$ proves the second.

By definition, G belongs to the union precisely when at least one $\mathbf{C} \in \mathfrak{T}_P(\mathcal{C})$ satisfies $G \in \mathcal{F}_P(\mathbf{C})$. This is equivalent to membership in $\mathcal{E}_P^*(G) \cap \mathfrak{T}_P(\mathcal{C})$. Finally, every component in an input from $\mathfrak{T}_P(\mathcal{C})$ belongs to its role-specific catalogue by construction, so component novelty is impossible. $\square$

E.4 PROGRAM-MASKED DISCRETE FLOW

For a program P , let $\mathcal{I}(P)$ index the categorical coordinates of its sampling layout. Define

$$\mathcal{S}(P) = \left\{ (q, j) \in \mathcal{I}(P) : m_j^{(q)}(P) = 0 \right\} \quad (33)$$

and let $x_{P,j}^{(q)}$ be the prescribed value of each fixed coordinate. The program-compatible face of the raw categorical state space is

$$\Omega_P = \left\{ x : x_j^{(q)} = x_{P,j}^{(q)} \text{ for every } (q, j) \in \mathcal{S}(P) \right\}. \quad (34)$$

Define the projection Π_P coordinatewise by

$$[\Pi_P(x)]_j^{(q)} = \begin{cases} x_{P,j}^{(q)}, & (q, j) \in \mathcal{S}(P), \\ x_j^{(q)}, & (q, j) \notin \mathcal{S}(P). \end{cases} \quad (35)$$

Training corruption has the form $X_t = \Pi_P(\tilde{X}_t)$. Numerical sampling begins from $X_0 = \Pi_P(\tilde{X}_0)$ and applies

$$X_{k+1} = \Pi_P(\Phi_k^\theta(X_k, \xi_k)), \quad k = 0, \dots, L-1, \quad (36)$$

where Φ_k^θ is the unconstrained stochastic update and ξ_k contains its sampling randomness.

Proposition E.4 (Program-fixed invariance). *For every program P ,*

$$\Pr[X_t \in \Omega_P] = 1 \quad (37)$$

at every corruption time t , and

$$\Pr[X_k \in \Omega_P \text{ for every } k = 0, \dots, L] = 1 \quad (38)$$

during numerical sampling.

Proof. By definition, $\Pi_P(x) \in \Omega_P$ for every raw state x , and $\Pi_P(x) = x$ for every $x \in \Omega_P$. Since corruption returns $\Pi_P(\tilde{X}_t)$, every corrupted state lies in Ω_P .

The sampling claim follows by induction. The projected initial state belongs to Ω_P . If $X_k \in \Omega_P$, the unconstrained update may change any coordinate, but applying Π_P restores every coordinate in $\mathcal{S}(P)$. Hence $X_{k+1} \in \Omega_P$. The conclusion holds for all sampling steps. $\square$

Define the joint population law

$$\nu(dG, dP, dt, dX_t) = \mu_{\text{bal}}(dG, dP) \mathcal{U}(dt) p_t(dX_t | X_G, P). \quad (39)$$

For a variable coordinate, let

$$\rho_{qj}(a | X_t, t, P) = \Pr_\nu[X_{G,j}^{(q)} = a | X_t, t, P]. \quad (40)$$

Theorem E.5 (Population optimality of masked denoising). *Assume an unrestricted conditional model class and $\lambda_q > 0$. For every coordinate satisfying $m_j^{(q)}(P) = 1$, every population minimizer of $\mathcal{L}_{\text{FORGE}}$ satisfies*

$$p_\theta^*(X_{G,j}^{(q)} = a | X_t, t, P) = \rho_{qj}(a | X_t, t, P) \quad (41)$$

for every categorical state a , ν -almost surely. Coordinates with $m_j^{(q)}(P) = 0$ contribute no loss and are instead governed by Proposition E.4.

Proof. Fix a family q , a variable coordinate j , and a context $C = (X_t, t, P)$ for which the regular conditional distribution is defined. Let $p(a) = p_\theta(X_{G,j}^{(q)} = a | C)$ and $\rho(a) = \rho_{qj}(a | C)$. The conditional expected loss is

$$-\sum_a \rho(a) \log p(a) = -\sum_a \rho(a) \log \rho(a) + \sum_a \rho(a) \log \frac{\rho(a)}{p(a)} \quad (42)$$

$$= H(\rho) + D_{\text{KL}}(\rho || p). \quad (43)$$

The entropy term is independent of p . By non-negativity of Kullback–Leibler divergence, the expression is minimized exactly when $p = \rho$.

For fixed P and q , the factor $\lambda_q / \sum_j m_j^{(q)}(P)$ is positive and independent of the predicted categorical distribution. It therefore does not change the pointwise minimizer. Integrating over contexts and repeating the argument for every variable coordinate proves the result. Coordinates with zero masks do not appear in the objective. $\square$

E.5 MATCHED-BUDGET POST-HOC FILTERING

Fix a reaction program P , and define

$$A_P(G) = \mathbb{1}[G \in \mathcal{G}_{L1}(P)]. \quad (44)$$

For a proposal distribution Q on the bounded graph space, let

$$Z_Q(P) = \Pr_{G \sim Q}[A_P(G) = 1] = \mathbb{E}_{G \sim Q}[A_P(G)] \quad (45)$$

denote the exact-L1 acceptance probability.

Proposition E.6 (Exact law of matched-budget post-hoc filtering). *Suppose $Z_Q(P) > 0$. Rejection filtering produces the conditional distribution*

$$Q_{\text{post}}(g \mid P) = \Pr_{G \sim Q}[G = g \mid A_P(G) = 1] = \frac{Q(g)A_P(g)}{Z_Q(P)}. \quad (46)$$

For B independent attempts $G_1, \dots, G_B \sim Q$, define

$$N_B = \sum_{b=1}^B A_P(G_b). \quad (47)$$

Then

$$N_B \sim \text{Binomial}(B, Z_Q(P)), \quad (48)$$

and consequently

$$\mathbb{E}[N_B] = BZ_Q(P), \quad (49)$$

$$\text{Var}(N_B) = BZ_Q(P)(1 - Z_Q(P)), \quad (50)$$

$$\Pr[N_B = 0] = (1 - Z_Q(P))^B. \quad (51)$$

If T_K is the number of independent attempts required to obtain K accepted products, then

$$\mathbb{E}[T_K] = \frac{K}{Z_Q(P)}. \quad (52)$$

Finally, for a program-conditioned proposal Q_P and an unconditioned proposal Q_0 ,

$$\mathbb{E}[N_B^P] - \mathbb{E}[N_B^0] = B(Z_{Q_P}(P) - Z_{Q_0}(P)). \quad (53)$$

Proof. For any graph g , Bayes' rule gives

$$\Pr[G = g \mid A_P(G) = 1] = \frac{\Pr[A_P(G) = 1 \mid G = g] \Pr[G = g]}{\Pr[A_P(G) = 1]} \quad (54)$$

$$= \frac{A_P(g)Q(g)}{Z_Q(P)}. \quad (55)$$

Since the proposals are independent, the indicators $A_P(G_1), \dots, A_P(G_B)$ are independent Bernoulli random variables with common success probability $Z_Q(P)$. Their sum is therefore binomial, which gives its expectation, variance, and zero-acceptance probability.

The waiting time T_K has the negative-binomial distribution with success probability $Z_Q(P)$, whose expectation is $K/Z_Q(P)$. Applying the binomial expectation separately to Q_P and Q_0 and subtracting proves the final identity. $\square$

The proposition compares sampling efficiency only. It does not imply that conditioning improves $Z_Q(P)$; the matched production experiment tests that empirical claim.

E.6 RECURSIVE ROUTE CONSTRUCTION

Lemma E.7 (Termination of bounded route assessment). *Fix an assessment contract $\mathcal{B}$ with finite maximum route depth $D_{\mathcal{B}}$. Assume that each evidence query terminates and returns either a non-expansion outcome or an expansion with finitely many reactants. Then $\mathcal{A}_{\mathcal{B}}(G, P)$ terminates with either a finite route tree or a typed failure state.*

Proof. We use induction on the remaining depth $d = D_{\mathcal{B}} - k$ at a node of depth k . When $d = 0$, the assessor performs the terminating evidence query and either accepts a terminal node, returns its typed failure, or rejects an expansion with a depth-exhaustion state. Suppose the claim holds whenever fewer than $d > 0$ levels remain. At the current node, the evidence query either returns a non-expansion outcome immediately or a finite reactant tuple. In the latter case, each child is assessed at depth $k + 1$ and therefore has $d - 1$ levels remaining. Every child assessment terminates with a finite result by the induction hypothesis, and a finite tuple of finite child trees forms a finite parent tree. Finite expansion, proposal, verifier, and time budgets can only stop this recursion earlier. Rejecting an ancestor identity additionally prevents a circular molecular dependency from being accepted, but finite depth alone establishes termination. $\square$

Theorem E.8 (Soundness of returned synthesis dossiers). *Fix $\mathcal{B}$ as in Lemma E.7. If $\mathcal{A}_{\mathcal{B}}(G, P) = T$, then T is finite and $\text{Complete}(G, T, P) = 1$. In particular, G passes exact L1 verification under P , every L2 reaction in T reconstructs its target exactly under its registered forward transform, and every leaf has a current accepted record identifying the same constitutional structure as supplier-listed and available. A candidate tree with $\text{Complete}(G, T, P) = 0$ cannot be returned as a synthesis dossier. The outcome $\mathcal{A}_{\mathcal{B}}(G, P) = \perp$ states only that no complete dossier was returned under $\mathcal{B}$. It does not prove that no route exists under a larger search budget, a different admitted reaction set, or a later supplier snapshot.*

Proof. Finiteness follows from Lemma E.7. It remains to establish the three factors in the completeness predicate. The assessor begins recursive route construction only after exact L1 decomposition and forward replay have succeeded, so a returned dossier satisfies $\text{L1}(G, P) = 1$.

We next use structural induction on the finite route tree. At a leaf, the assessor returns a complete node only when the evidence policy admits an exact terminal-material record whose availability state is current and closed. Therefore $\text{Available}(\ell) = 1$ for every leaf ℓ of a returned tree. At an internal node, the assessor admits an expansion only when the registered reaction has exact-source, exact-substrate evidence and its deterministic forward transform uniquely reconstructs the target product. Thus $\text{Verify}(r) = 1$ for the associated reaction node r . The internal node is complete only when every child is complete; the induction hypothesis therefore applies to all descendant reactions and leaves.

All factors in $\text{Complete}(G, T, P)$ equal one, so their product equals one. Conversely, if any factor for a candidate tree equals zero, the fail-closed return rule does not admit that tree as a dossier and instead records the corresponding typed failure. Because the search and evidence contract is bounded, this rejection establishes no completeness claim about routes outside $\mathcal{B}$. $\square$

F GENERATED-STRUCTURE ATLAS

Table 12 reports 12 deterministic display examples from the same frozen seed-0 Ugi population used in Figure 2. The display begins with the lowest constitutional SHA-256 rank and then applies ECFP4 MaxMin selection among the 23 eligible, constitutionally unique products. This promotes visual coverage without using activity, route closure, procurement evidence or human preference. Each row includes the canonical constitutional SMILES so that the displayed identity is explicit; the exact strings are also stored in the hash-pinned atlas receipt.

Table 12: **FORGE generated-structure atlas**. The 12 rows form one multipage table, indexed by display order. Each constitutional graph is followed by its canonical SMILES. The decorated three-dimensional column shows one deterministic post-hoc RDKit ETKDGv3 structure optimized with up to 400 MMFF94s iterations. Purple denotes the amine-derived head, teal the aldehyde-derived body and tail, orange the isocyanide-derived tail and gray the assembly-introduced oxygen; these assignments are recovered through exact atom-mapped forward replay. FORGE does not generate three-dimensional coordinates, and these views are not predicted bioactive conformations. Exact L1 replay and descriptor-manifold membership are structural diagnostics, not evidence of activity, route closure or synthesis success. These are deterministic display examples, not experimental candidates.







CONSTITUTIONAL GRAPH AND CANONICAL SMILES		DECORATED 3D VIEW
1	 <chem>CCCCC(=O)CCCCCCCCCCCC(C(=O)NCCCOC)N(C)CCCC</chem>	
2	 <chem>CCCCCCCCCCCCCCCC(NCCCCCCC)C(=O)NCCCCSC(=O)CC1(C)OC1C</chem>	
3	 <chem>CC#CCCCCCCCC=CCCCC(NCCCC(=O)CC(C)(C)C)C(=O)NCCCCC</chem>	

Table 12, continued


















	CONSTITUTIONAL GRAPH AND CANONICAL SMILES	DECORATED 3D VIEW
4	 <chem>CCCCCCCCCCCCC(=O)NC1CCCC1)N(C)CC(C)CC</chem>	
5	 <chem>C#CCC(CC)CCCC=CC#CCCCCCCCCCCC(=O)CCCC(NCCC(C)CCCC)C(=O)NCCCOC</chem>	
6	 <chem>CCCCCCCCCCCC=CCCC(C(=O)NCCC1CCCC1)N(C)CC(C)(C)C</chem>	
7	 <chem>CCCCC(=O)CCCCCCCCCCCC(NCCC(C)(C)C)C(=O)NCC1CO1</chem>	
8	 <chem>CCCCC(C)CC(C(=O)NCCC1OCC1)N(C)CC(C)CC</chem>	

Table 12, continued

	CONSTITUTIONAL GRAPH AND CANONICAL SMILES	DECORATED 3D VIEW
9	 <chem>CCCCCCCCC=CCCCC(C(=O)NCCCCCOC1CCOC1)N(C)CCCC(C)CC</chem>	
10	 <chem>CCCCCCCCC(NCCC(C)(C)C)C(=O)NCCCCCOC=O</chem>	
11	 <chem>CCCCCNC(CCCCCCCCCC(=O)CCC)C(=O)NC1CCO1</chem>	
12	 <chem>CCCCCCCCC(=O)CCCCCCC(C(=O)NCOCOC)N(C)CC(C)CC</chem>	